

Graph Neural Networks: Development and Deployment in Particle Physics

Lazar Novakovic and Loukas Gouskos

Brown University, Department of Physics

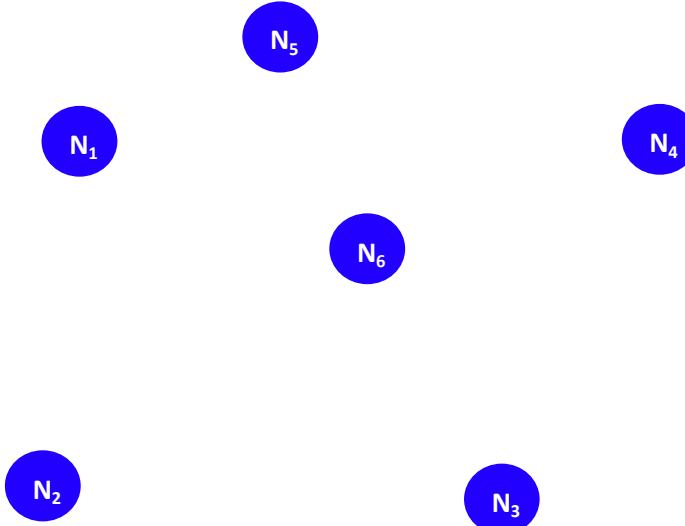
AI Winter AI School (2025)

What is a Graph?

- A mathematical structure used to model pairwise relations between points (objects/particles/measurements)

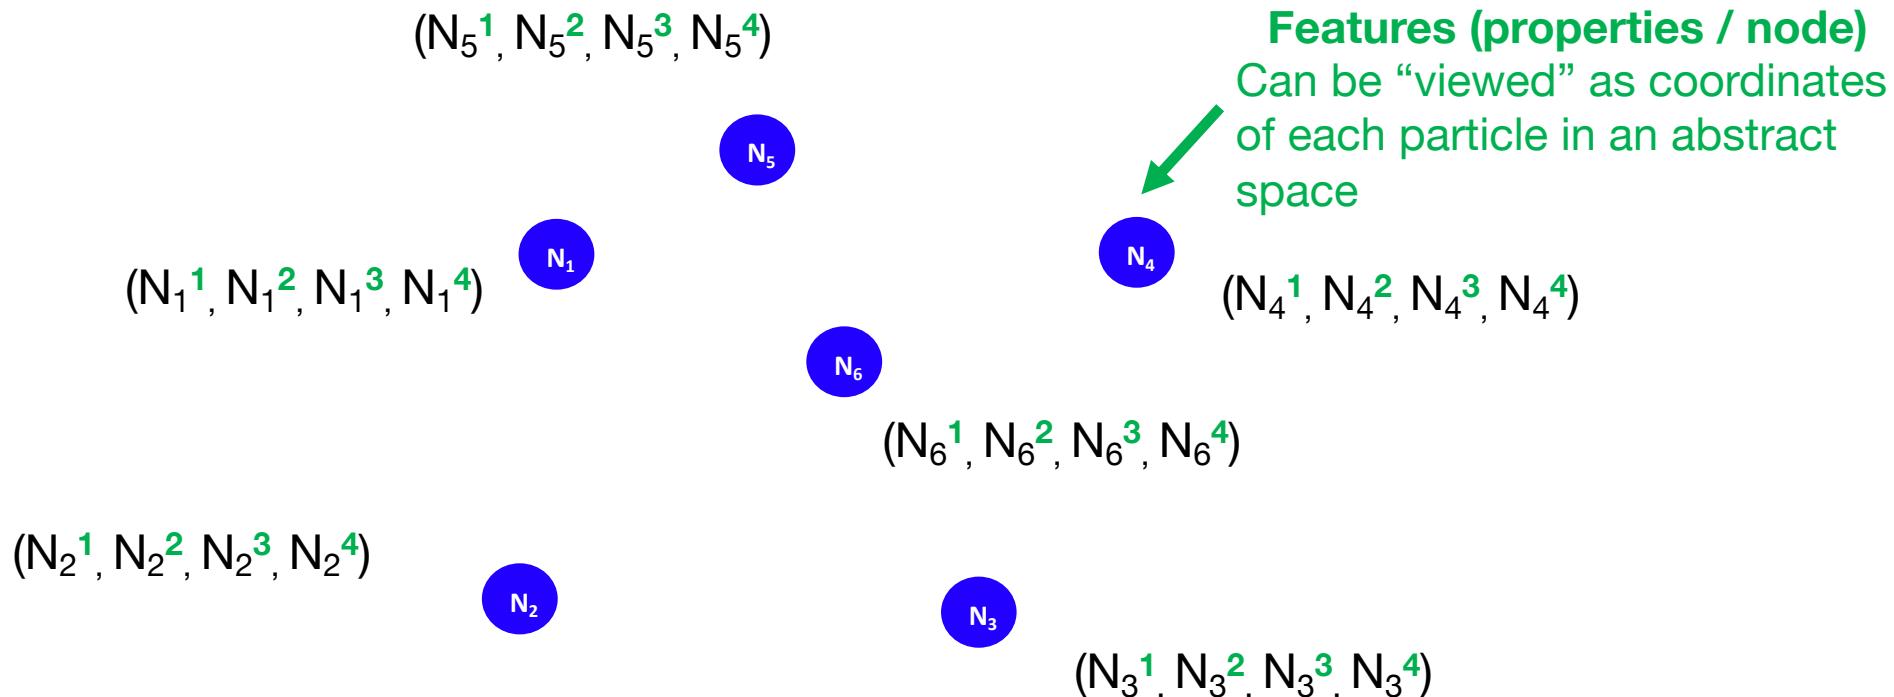
What is a Graph?

- A mathematical structure used to model pairwise relations between points (objects/particles/measurements)
 - ◆ Set of **nodes/vertices (N_i)**



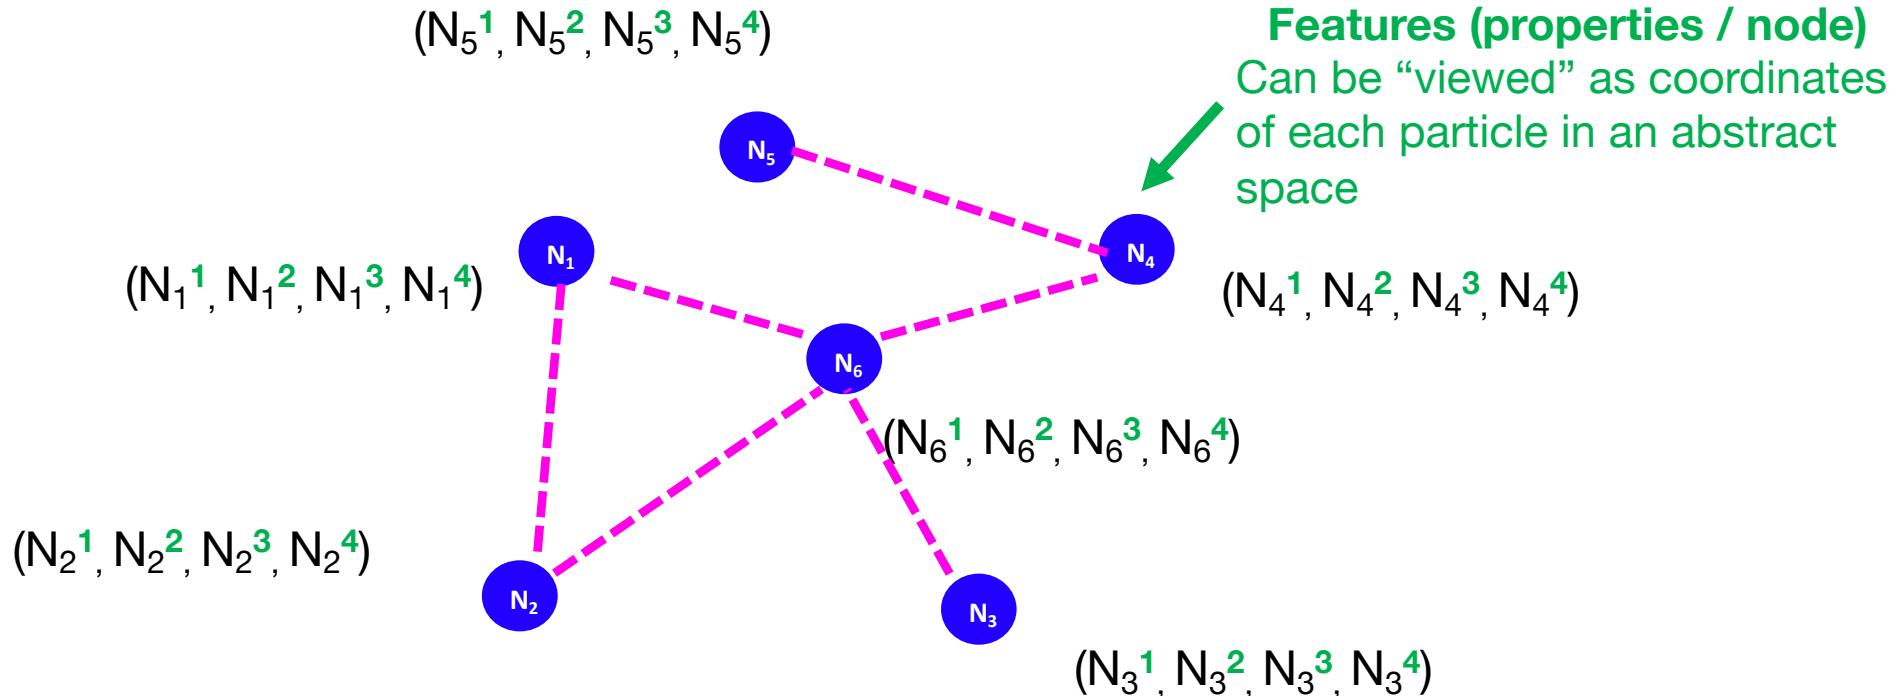
What is a Graph?

- A mathematical structure used to model pairwise relations between points (objects/particles/measurements)
 - ◆ Set of **nodes/vertices (N_i)**



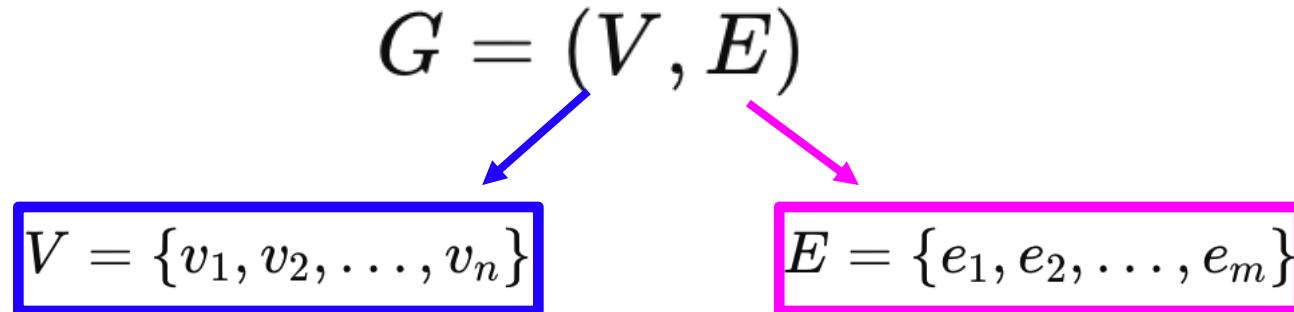
What is a Graph?

- A mathematical structure used to model pairwise relations between points (objects/particles/measurements)
 - ◆ Set of **nodes/vertices (N_i)**
 - ◆ Set of **edges/connections** between **nodes**



What is a Graph?

- A mathematical structure used to model pairwise relations between points (objects/particles/measurements)
 - ◆ Set of **nodes/vertices (N ,**
 - ◆ Set of **edges/connections** between **nodes**



What is a Graph?

- Mathematical representation via **adjacency matrix** (A)

- ◆ Vertices: $V = \{v_1, v_2, v_3\}$

- ◆ Edges: $E = \{\{v_1, v_2\}, \{v_2, v_3\}\}$

- ◆ If vertices i,j form an edge:

$$a_{ij} = 1$$

- or equal to the “weight” of the edge

- ◆ else:

$$a_{ij} = 0$$

What is a Graph?

- Mathematical representation via **adjacency matrix** (A)

- ◆ Vertices: $V = \{v_1, v_2, v_3\}$

- ◆ Edges: $E = \{\{v_1, v_2\}, \{v_2, v_3\}\}$

- ◆ If vertices i,j form an edge:

$$a_{ij} = 1$$

$$A = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix}$$

- or equal to the “weight” of the edge

- ◆ else:

$$a_{ij} = 0$$

Graph Neural Networks

- Graphs: Powerful math structures to model relationships
 - ◆ Yet, simply representing data as a graph isn't enough

Graph Neural Networks

- Graphs: Powerful math structures to model relationships
 - ◆ Yet, simply representing data as a graph isn't enough
- **Graph Neural Networks (GNNs)** come into play!

Graph Neural Networks

- Graphs: Powerful math structures to model relationships
 - ◆ Yet, simply representing data as a graph isn't enough
- **Graph Neural Networks (GNNs) come into play!**
 - ◆ Real-world data: complex relationships
 - Traditional ML: struggle with non-Euclidean data
 - ◆ GNNs: can learn to extract features from graph structures
 - ideal where relationships (edges) b/w entities (nodes) matter

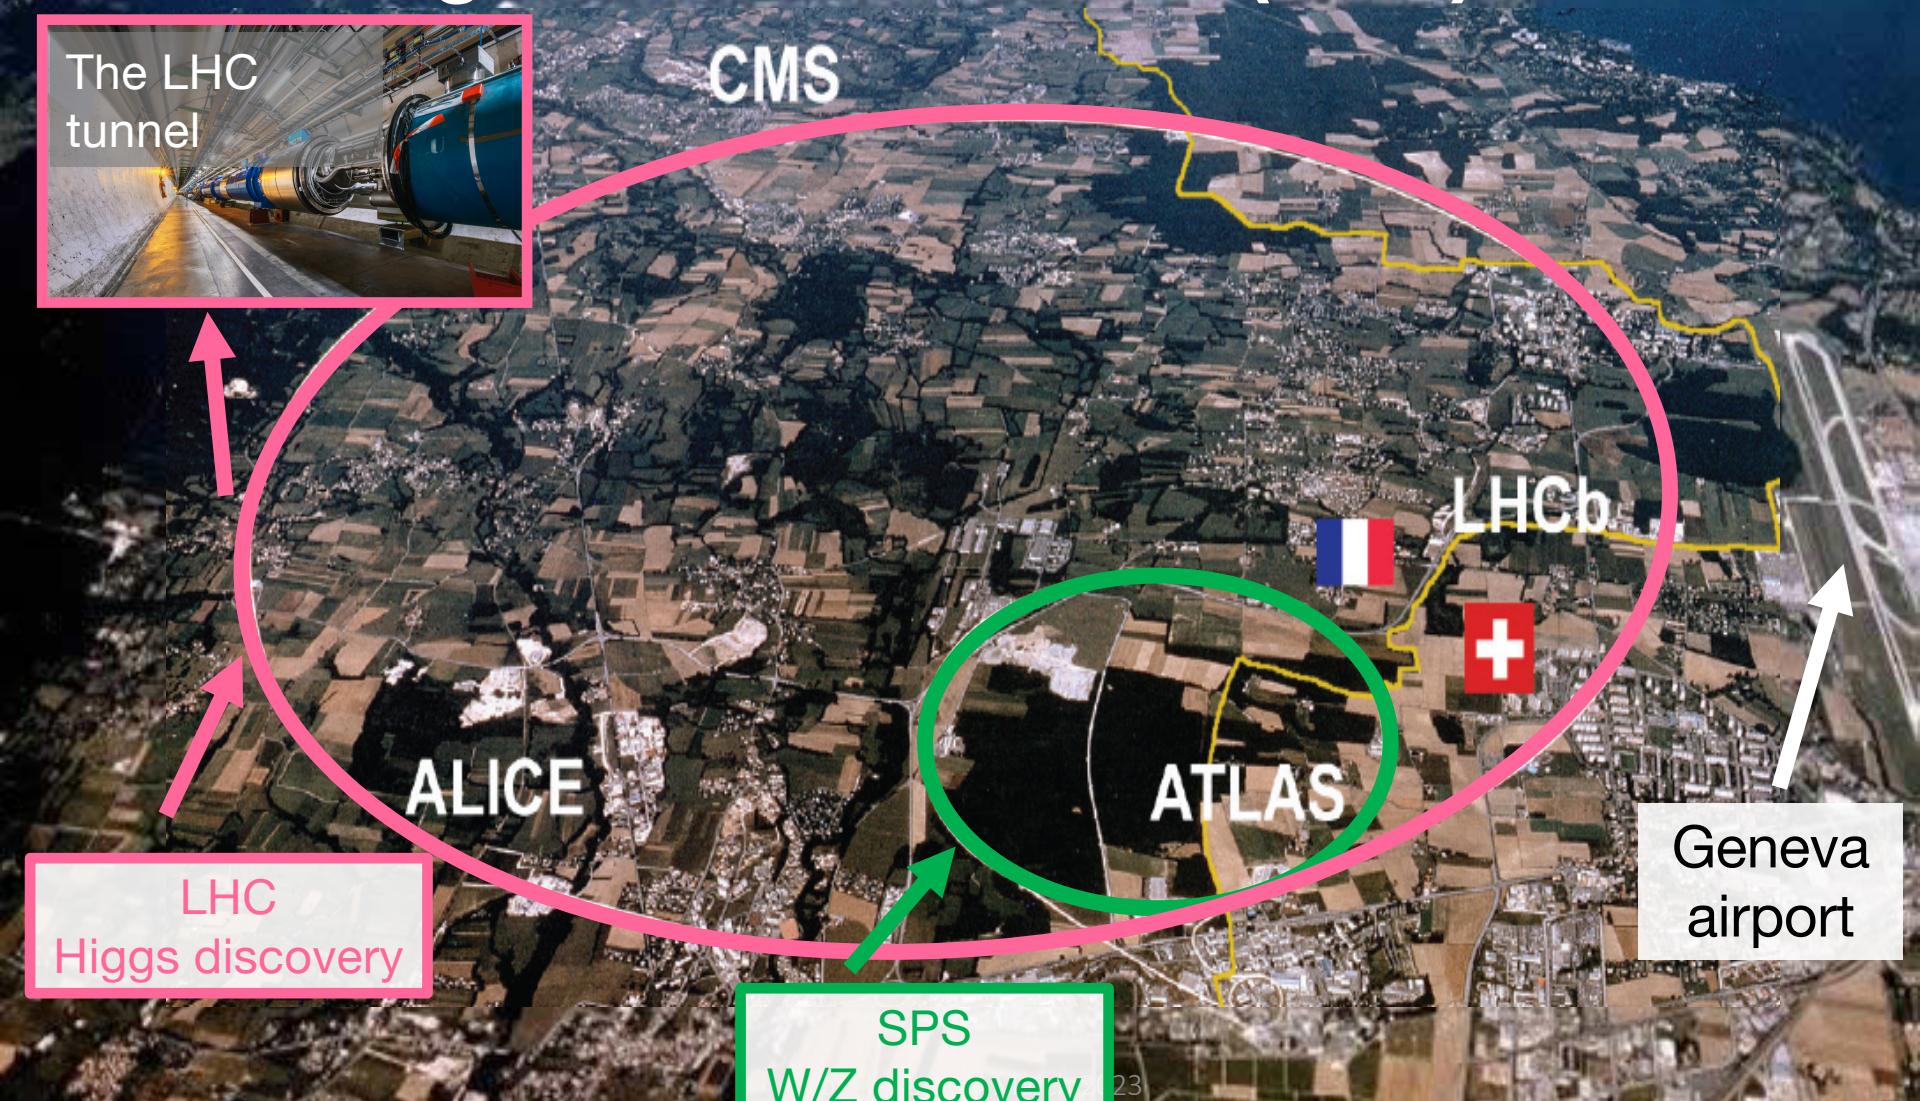
Graph Neural Networks

- Graphs: Powerful math structures to model relationships
 - ◆ Yet, simply representing data as a graph isn't enough
- **Graph Neural Networks (GNNs) come into play!**
 - ◆ Real-world data: complex relationships
 - Traditional ML: struggle with non-Euclidean data
 - ◆ GNNs: can learn to extract features from graph structures
 - ideal where relationships (edges) b/w entities (nodes) matter
- Core Idea:
 - ◆ Message Passing
 - Nodes exchange info w/ their neighbors to aggregate local context
 - ◆ Feature Transformation
 - Node features updated based on neighboring info (e.g., via NN)
 - ◆ Iterative Updates (repeat over multiple layers)
 - Info flows across Graph → Capture long-range dependences

Overview of this lecture

- Executive summary: LHC experiments
- Why Graph Neural Networks (GNNs) in Particle Physics
- Hands-on:
 - ◆ GNNs for classification (i.e., discriminate b/w different particles)
 - ◆ GNNs for clustering (i.e., group detector signals)

The Large Hadron Collider (LHC) at CERN



The Large Hadron Collider (LHC) at CERN

The LHC
tunnel



CMS

- **LHC:** The largest particle collider in the world

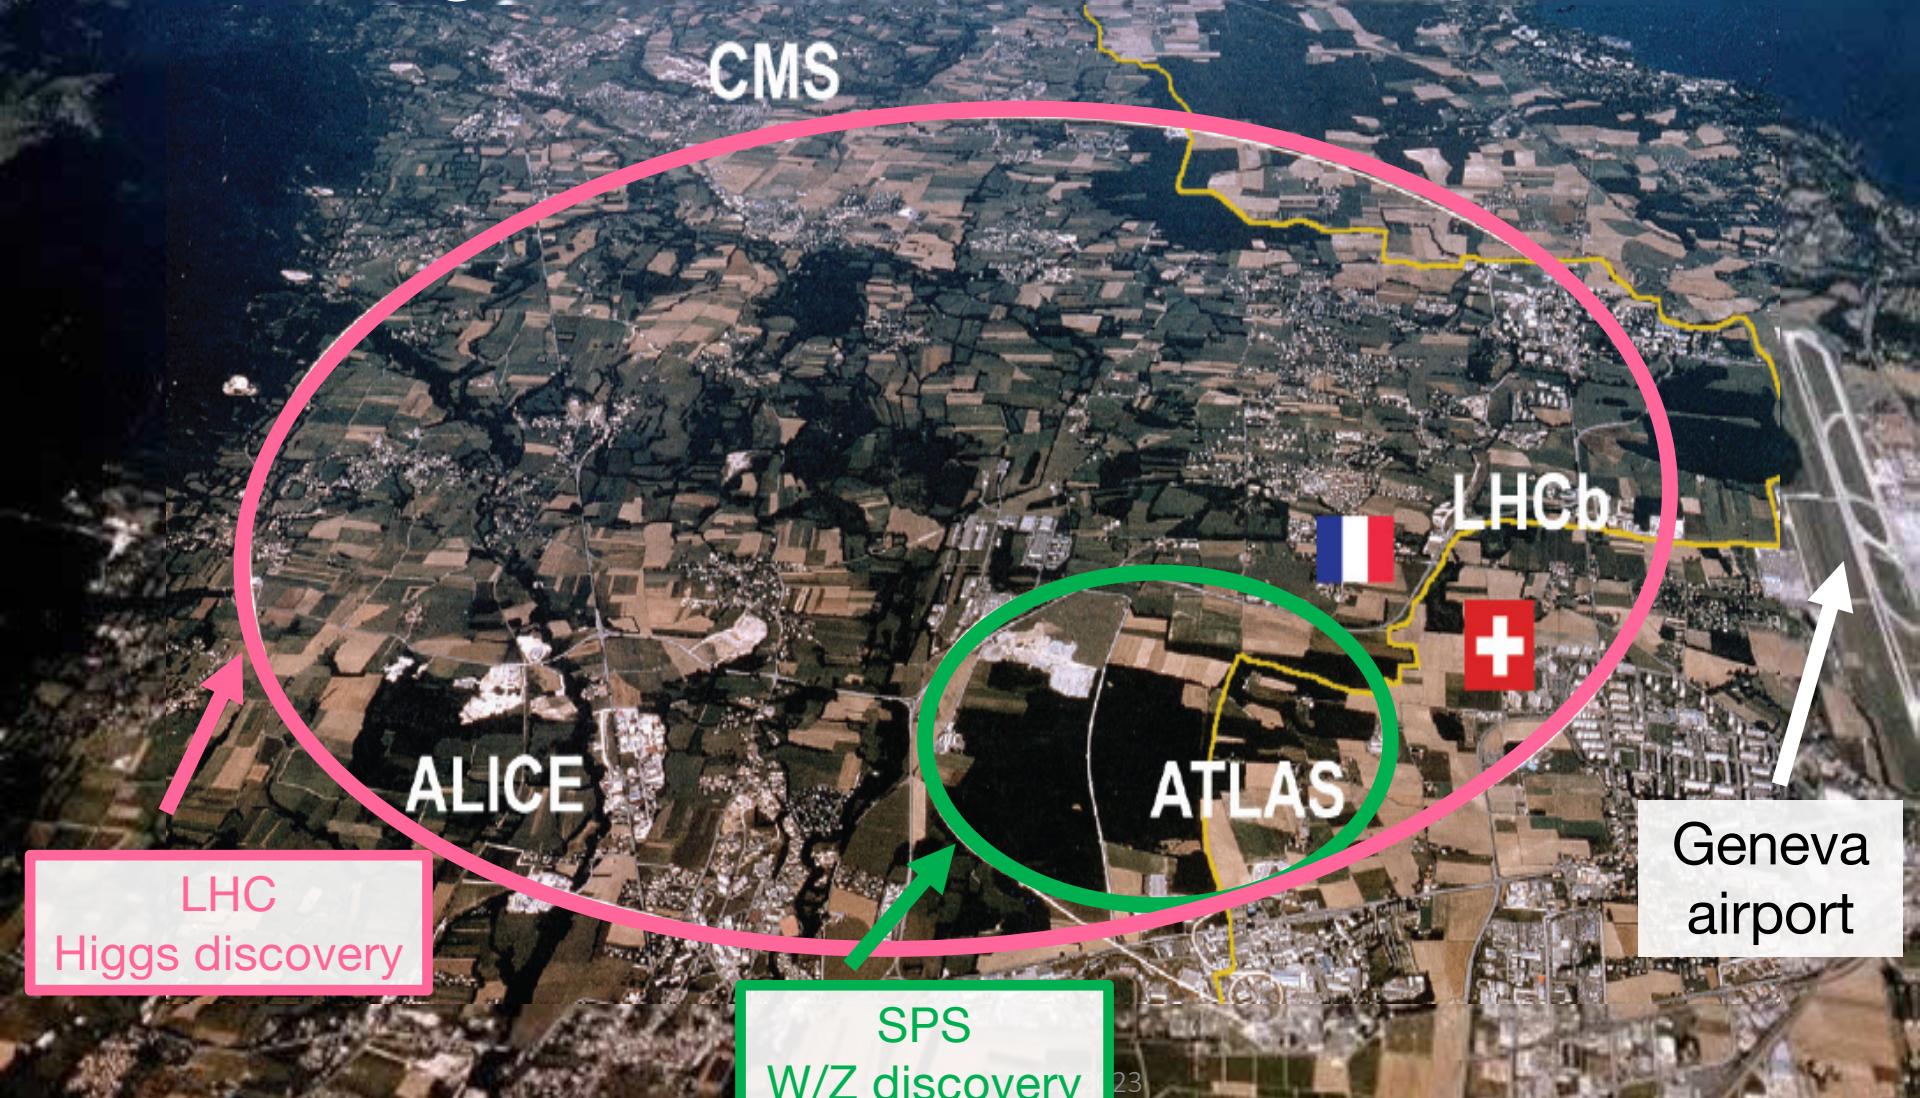
- ◆ **27km** perimeter
- ◆ >1.2K dipole magnets @ 8T
 - cooled @ -271° C
- ◆ Colliding protons at energies: **7 – 14 TeV**
 - currently @13.6 TeV (Run 3)
[almost 14K times proton mass]
- ◆ 2808 bunches collide every **25nsec**
 $O(10^{12})$ protons/bunch

LHC
Higgs discovery

Geneva
airport

SPS
W/Z discovery

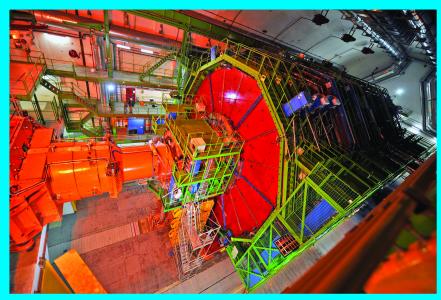
The Large Hadron Collider (LHC) at CERN



LHC
Higgs discovery

SPS
W/Z discovery

The Large Hadron Collider (LHC) at CERN

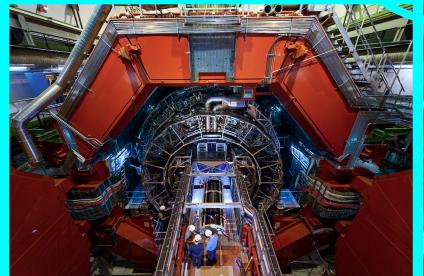


CMS



LHCb

ALICE



ATLAS



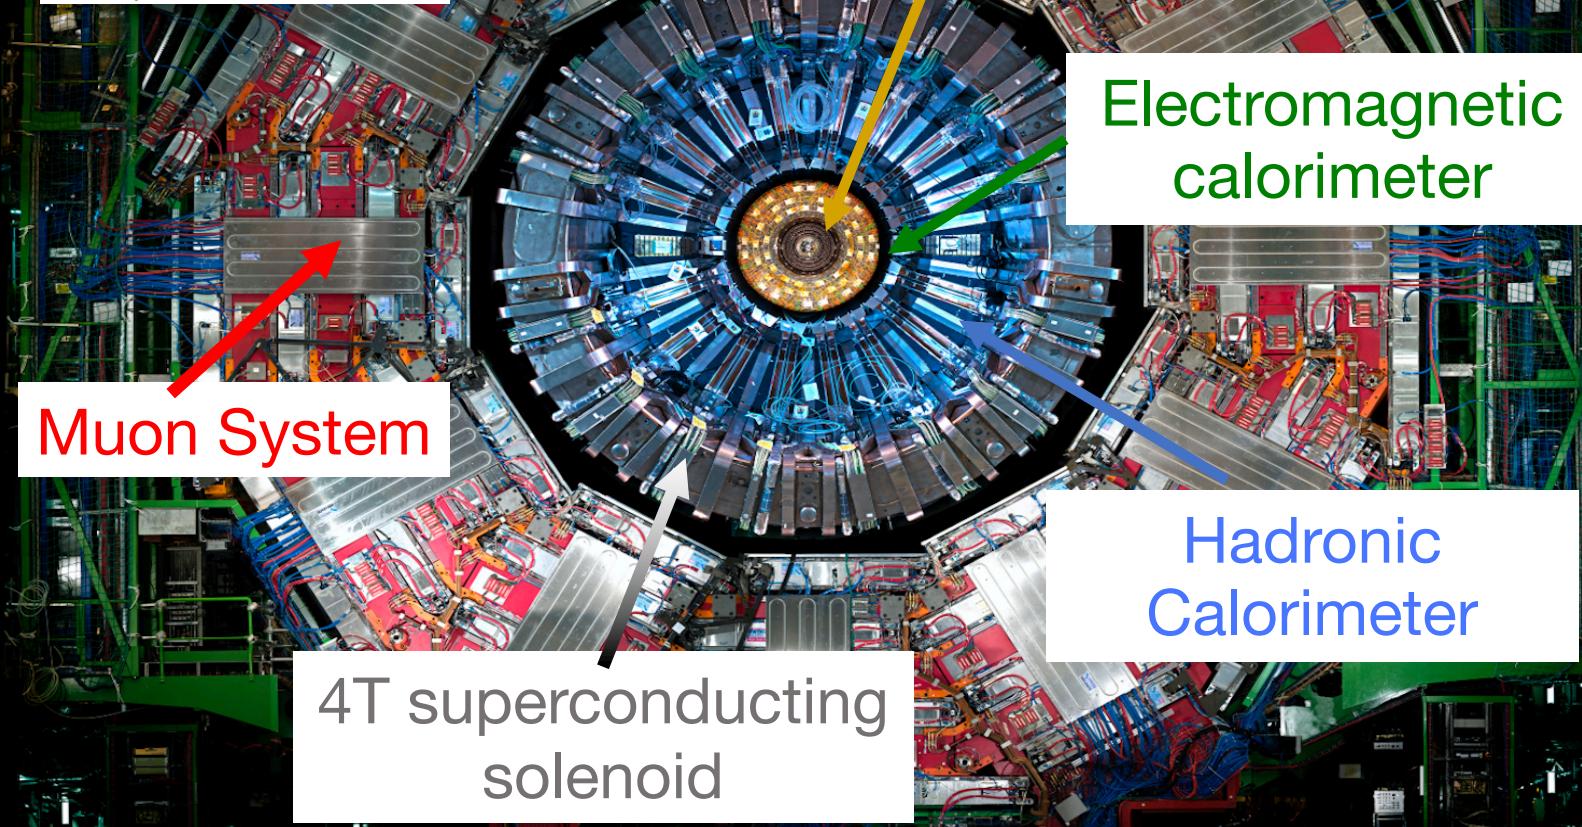
The LHC detectors [e.g., CMS]

CMS "cheat sheet"

Weight : 14.000 tons

Diameter: ~ 15m

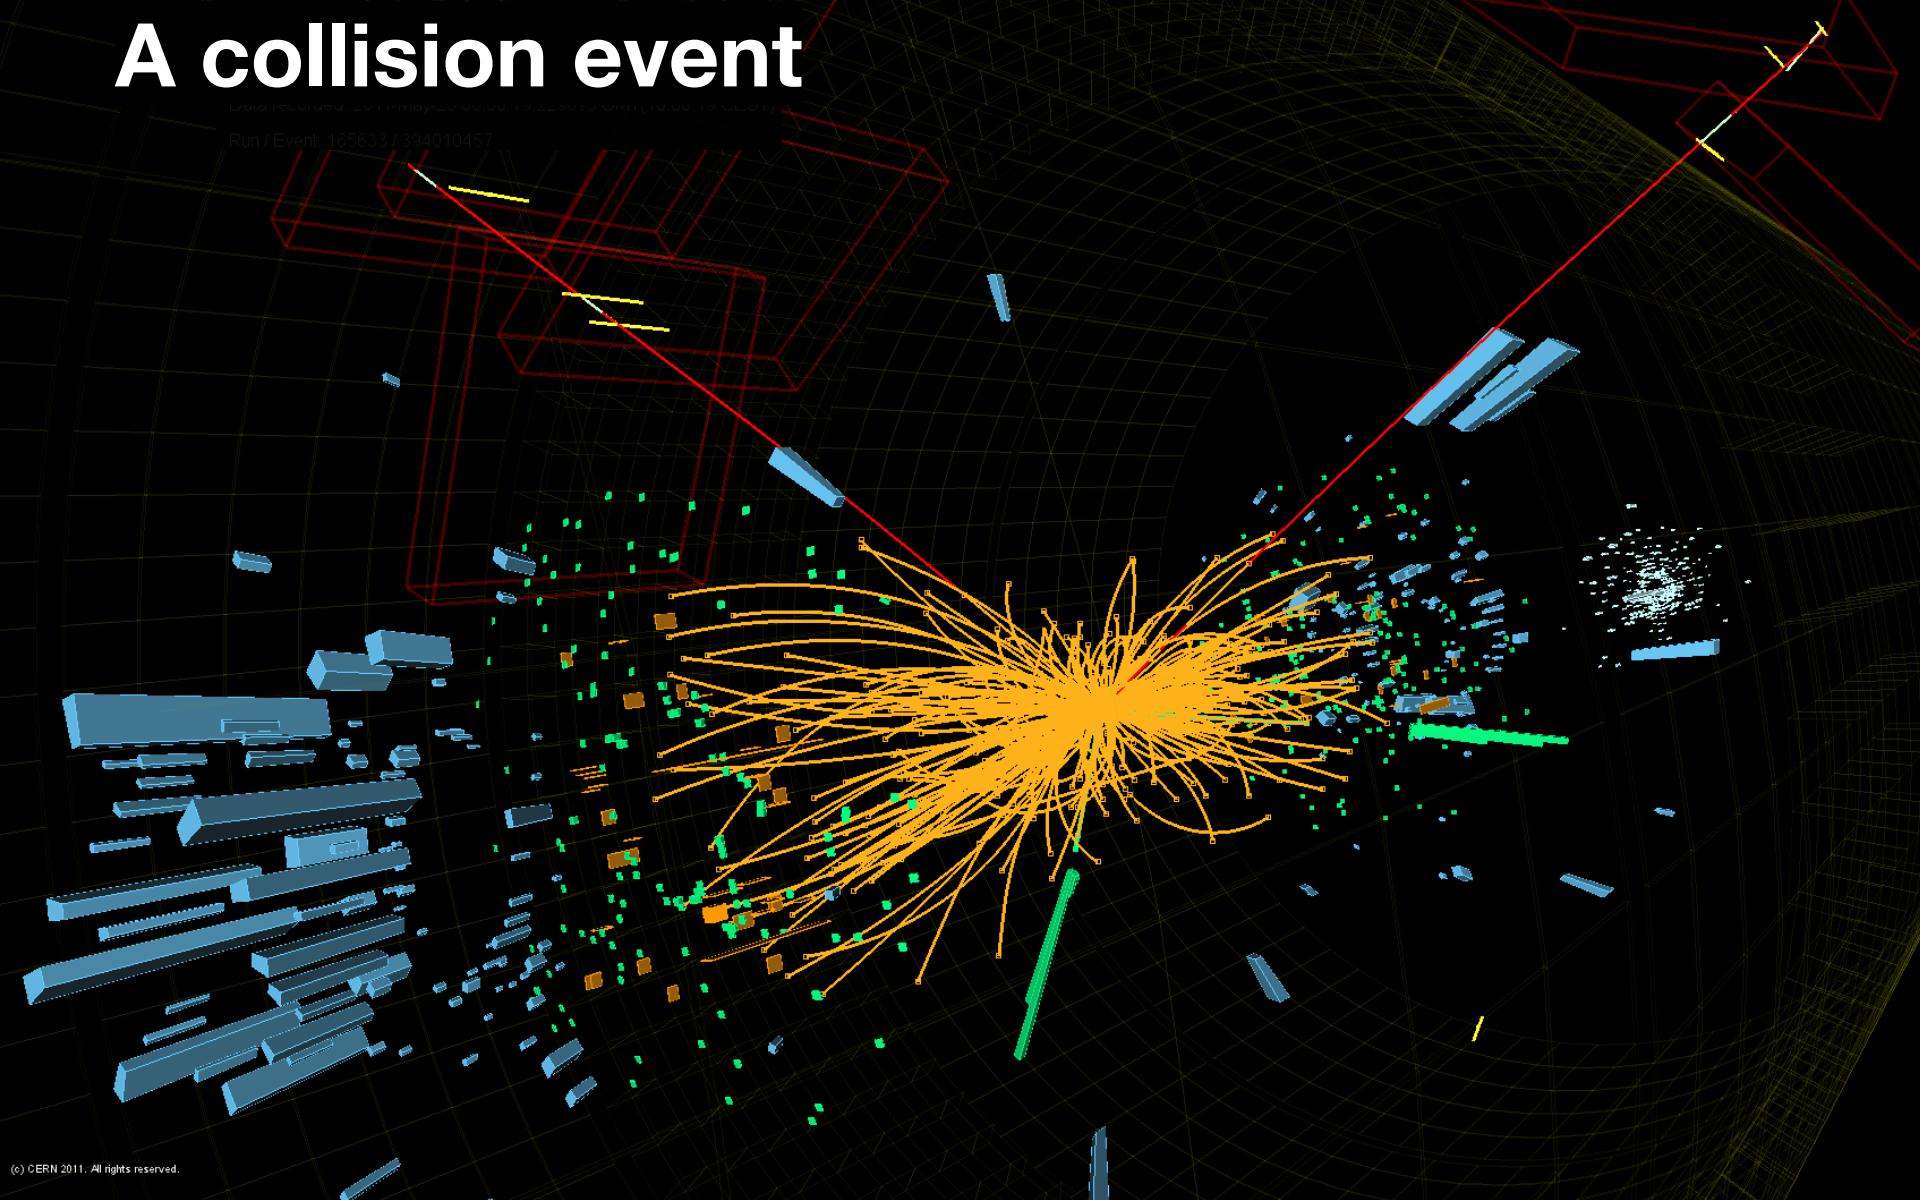
Length : ~ 23 m



A collision event

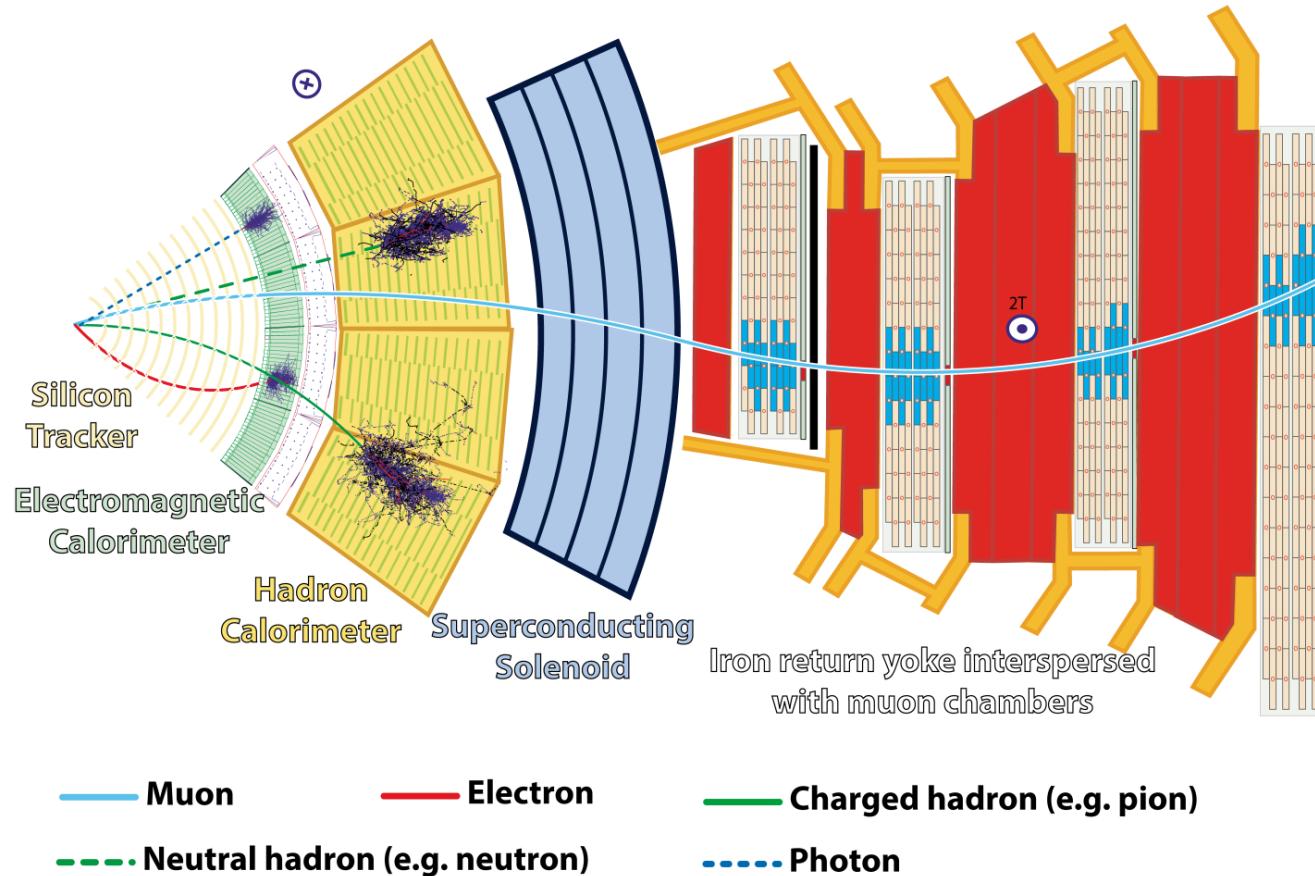
Date received: 2011-09-20 00:00:12.229730 GMT | 19:00:12 UTC 2011

Run / Event: 165633 / 394010457

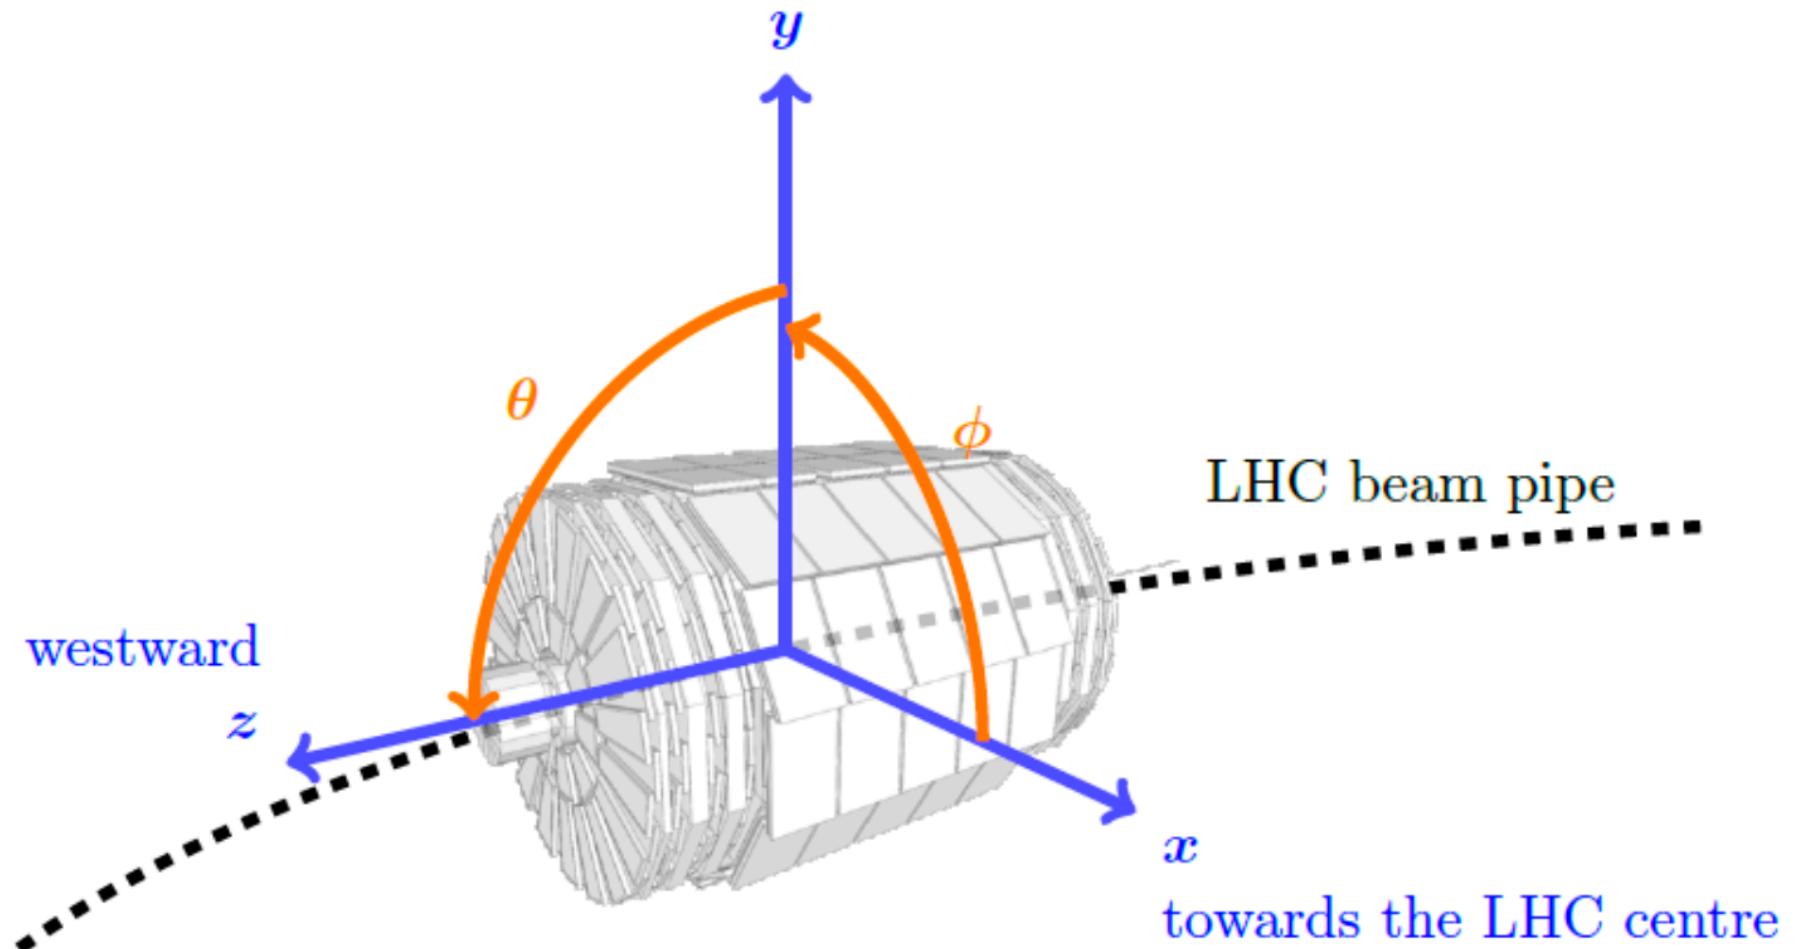


Event reconstruction

- Combine information from all subdetectors
 - ◆ Mutually exclusive list of particles [i.e. Particle Flow]



towards the surface



towards the LHC centre

Key points

- **High-dimensional** and **sparse** data
 - ◆ Complex events, high-dimensional data, where most regions contain little or no signal → sparse data
- **Heterogeneous** events
 - ◆ different types of particles (e.g., muons, photons, hadrons)
 - ◆ each interacting differently with the sub-detectors
- **Long-range** dependencies
 - ◆ Physical processes often involve relationships between distant regions of the detector
 - e.g., correlations between particles traveling in different directions
- **Complex** topology
 - ◆ Particle collisions produce jets, tracks, and vertices that naturally form graph-like structures rather than regular grids

Advantages of GNNs

- Flexibility

- ◆ Do not impose a fixed structure on the data
 - highly adaptable to various structures
 - regular grids (where they can approximate CNNs) to highly irregular networks

Advantages of GNNs

- Flexibility

- ◆ Do not impose a fixed structure on the data
 - highly adaptable to various structures
 - regular grids (where they can approximate CNNs) to highly irregular networks

- Local patterns

- ◆ Exploit local structure in graphs [such as CNNs].
 - GNNs: capture the properties of a node based on its surroundings, respecting the graph topology

Advantages of GNNs

- Flexibility

- ◆ Do not impose a fixed structure on the data
 - highly adaptable to various structures
 - regular grids (where they can approximate CNNs) to highly irregular networks

- Local patterns

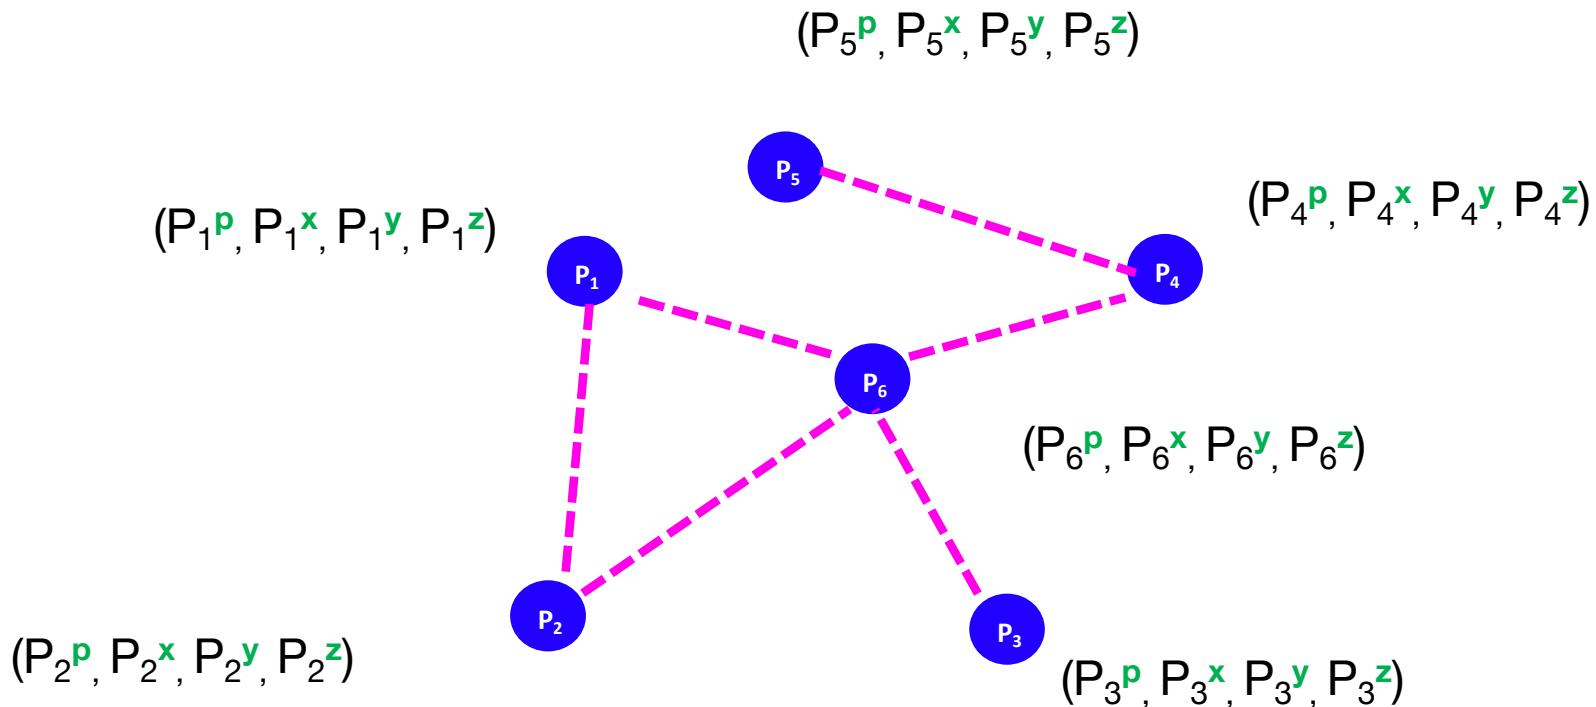
- ◆ Exploit local structure in graphs [such as CNNs].
 - GNNs: capture the properties of a node based on its surroundings, respecting the graph topology

- Efficiency:

- ◆ More Parameter-efficient than dense networks
 - GNNs focus on “relationships” between input data that matter
 - ignore the irrelevant ones

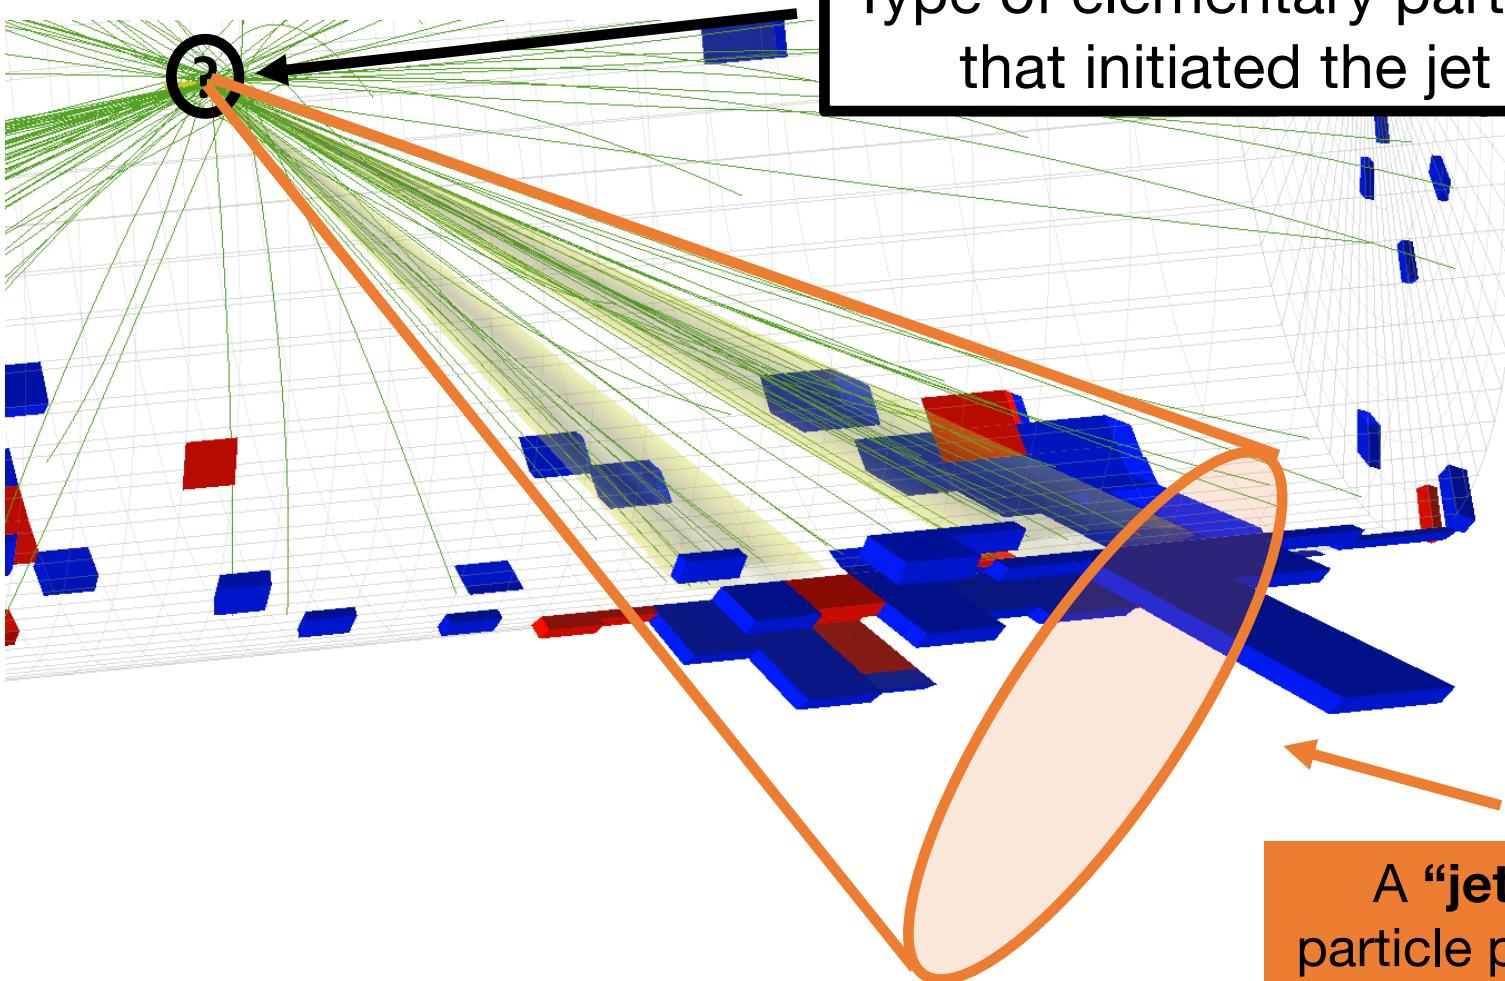
GNNs in particle physics

- **Nodes:** particles, detector signals, etc..
- **Features:** momentum, spatial coo/ntes (x,y,z), ID, ..
- **Edges:** correlations between the particles

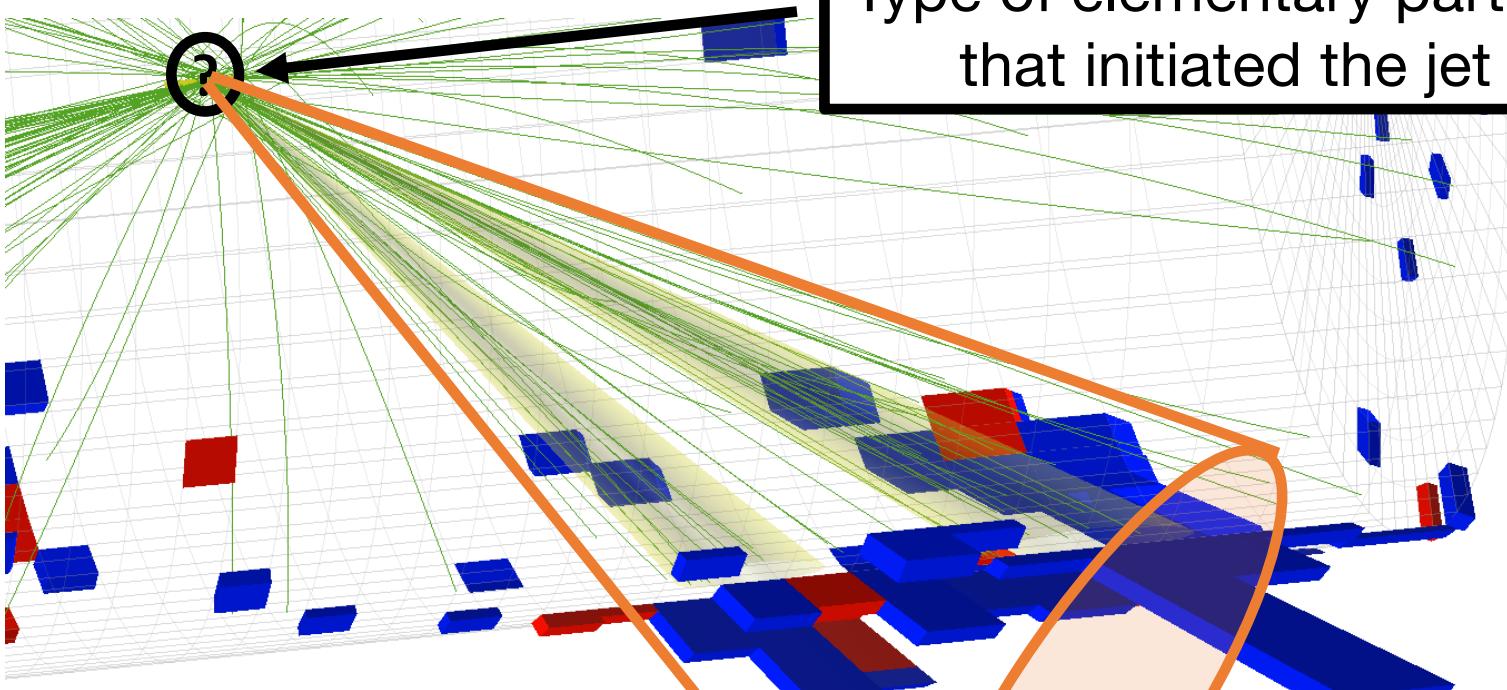


Exercise 1: Classify different particle types

Exercise 1: Classify different particle types



Exercise 1: Classify different particle types



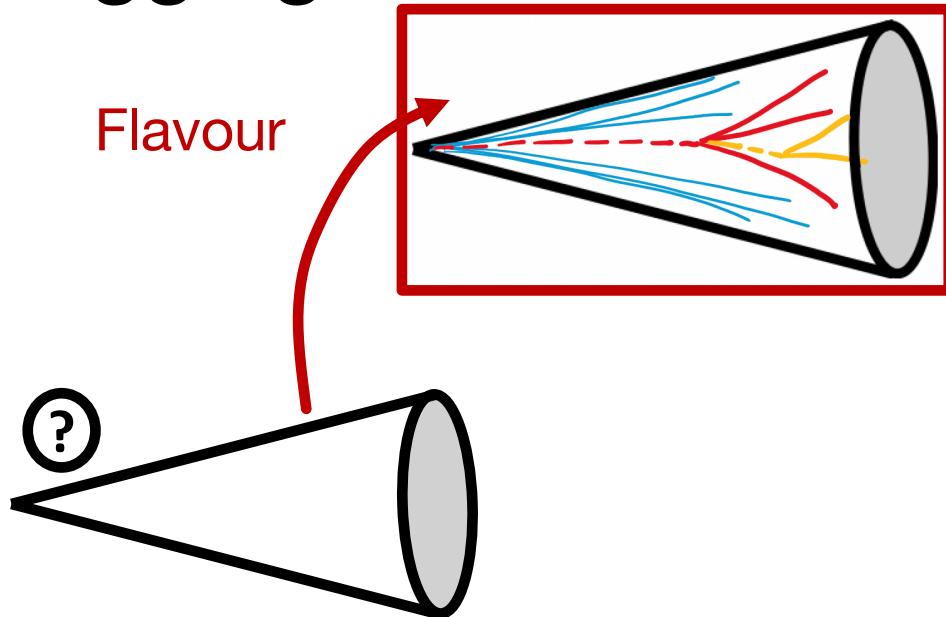
Type of elementary particle
that initiated the jet

Facts:

- $O(50-100)$ particles per jet
- $O(50)$ features per particle
- $\sim O(1000)$ inputs/jet

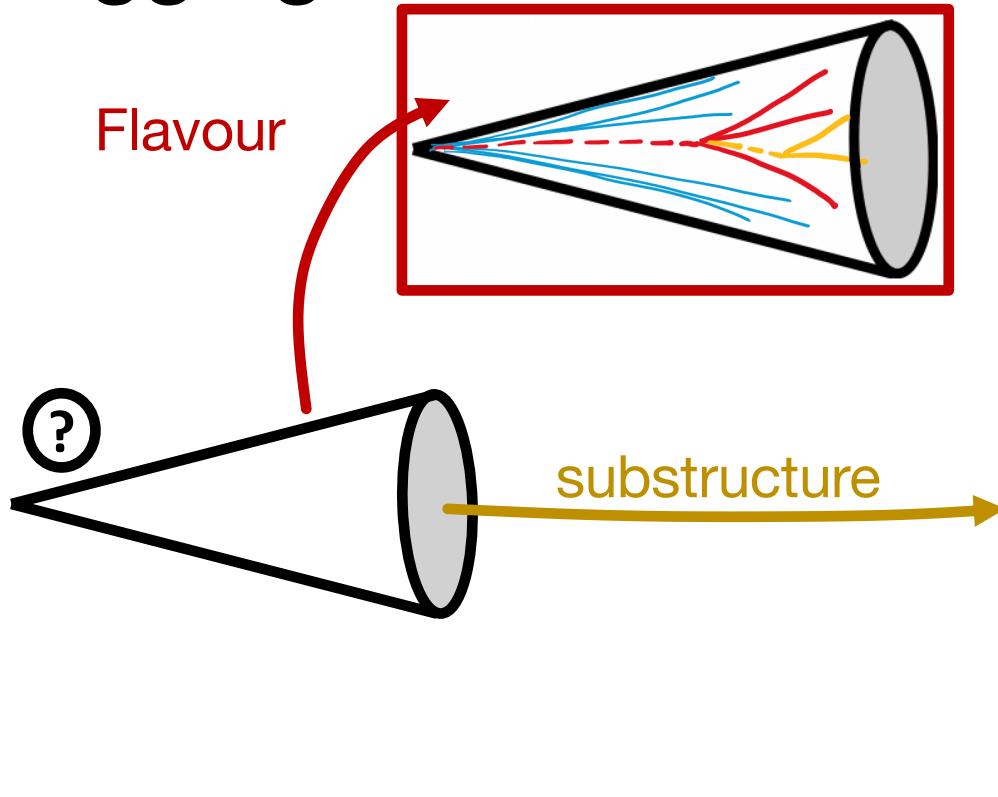
A “jet” in
particle physics

Jet tagging basics

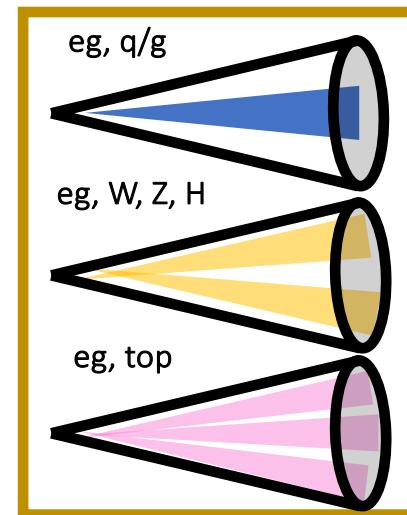


- Larger lifetime of b/c hadrons
- Displaced tracks/vertices
- Harder fragmentation
- Non-isolated leptons

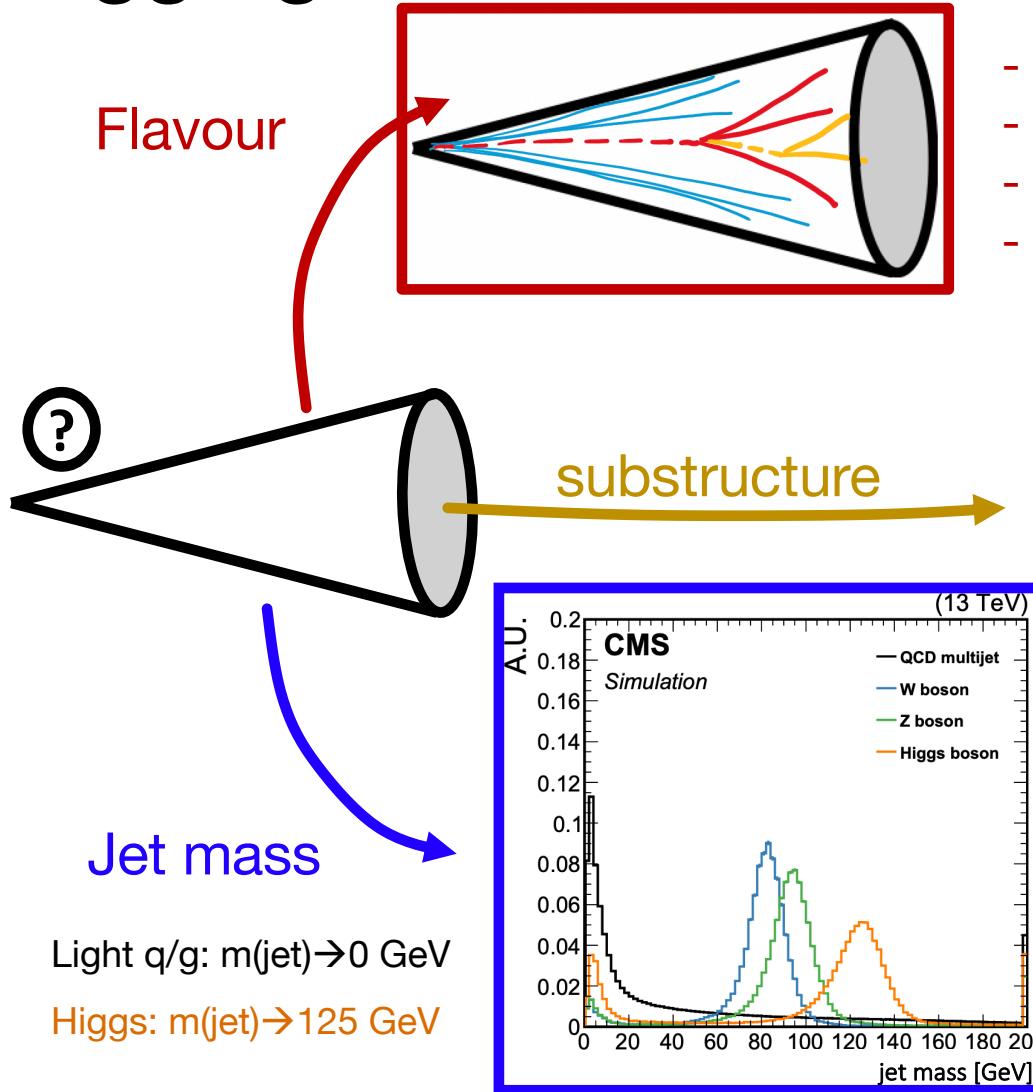
Jet tagging basics



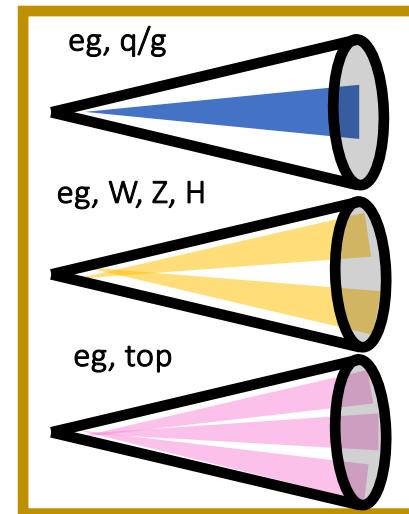
- Larger lifetime of b/c hadrons
- Displaced tracks/vertices
- Harder fragmentation
- Non-isolated leptons



Jet tagging basics



- Larger lifetime of b/c hadrons
- Displaced tracks/vertices
- Harder fragmentation
- Non-isolated leptons



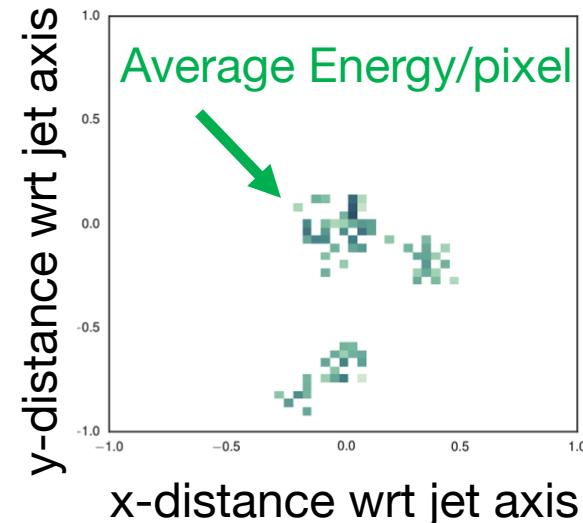
How to “represent” the jet

How to “represent” the jet

- Key for powerful AI/ML-based algorithms
 - ◆ We are going to compare two representations
 - Jet as an Image: Process via Convolutional Neural Networks (CNNs)
 - Jet as a Graph: Process via GNNs
 - ◆ Compare them

Jets as images

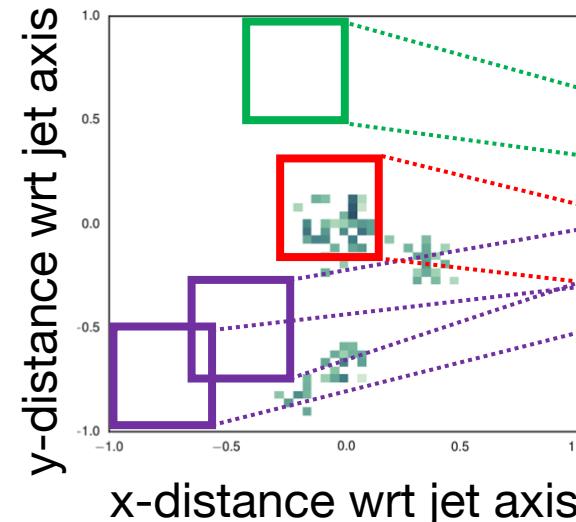
Jet Image
[a top quark]



detector → camera
jet → image

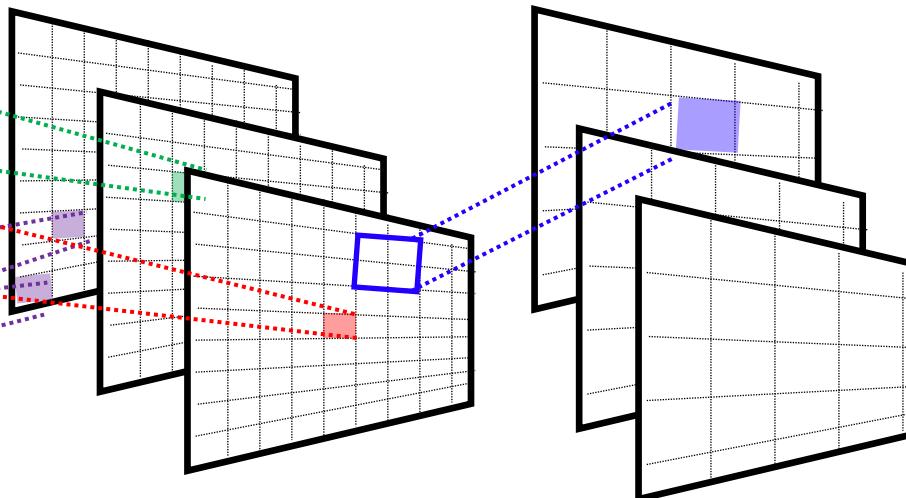
Jets as images

Jet Image
[a top quark]



detector → camera
jet → image

Extract features



2D-conv layers

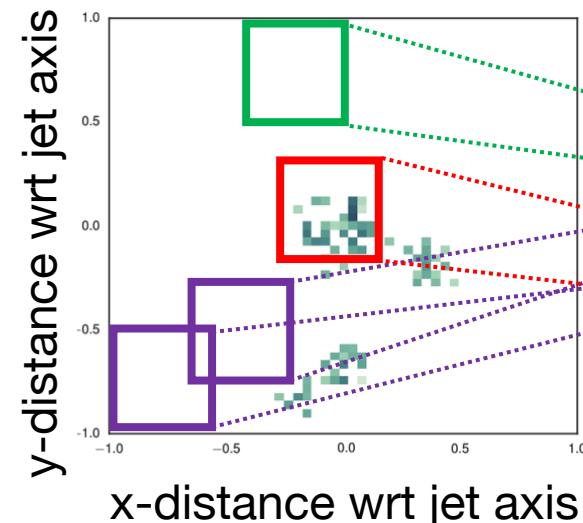
- Convolution layers can be stacked
- Deeper layers exploit more general features

Pooling

- Max, Min,... of an area [e.g., 2x2 pix] of the feature map
- Robustness
- Dimensionality reduction

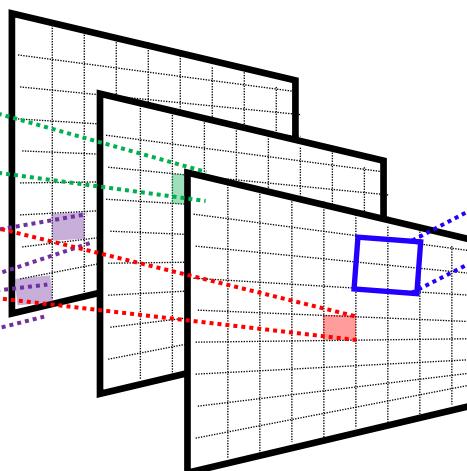
Jets as images

Jet Image [a top quark]



detector → camera
jet → image

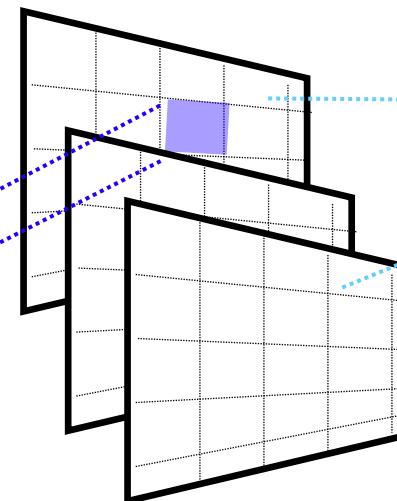
Extract features



2D-conv layers

- Convolution layers can be stacked
- Deeper layers exploit more general features

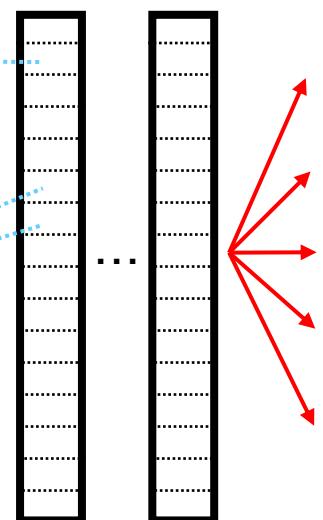
Extract features



Pooling

- Max, Min,... of an area [e.g., 2x2 pix] of the feature map
- Robustness
- Dimensionality reduction

Classification



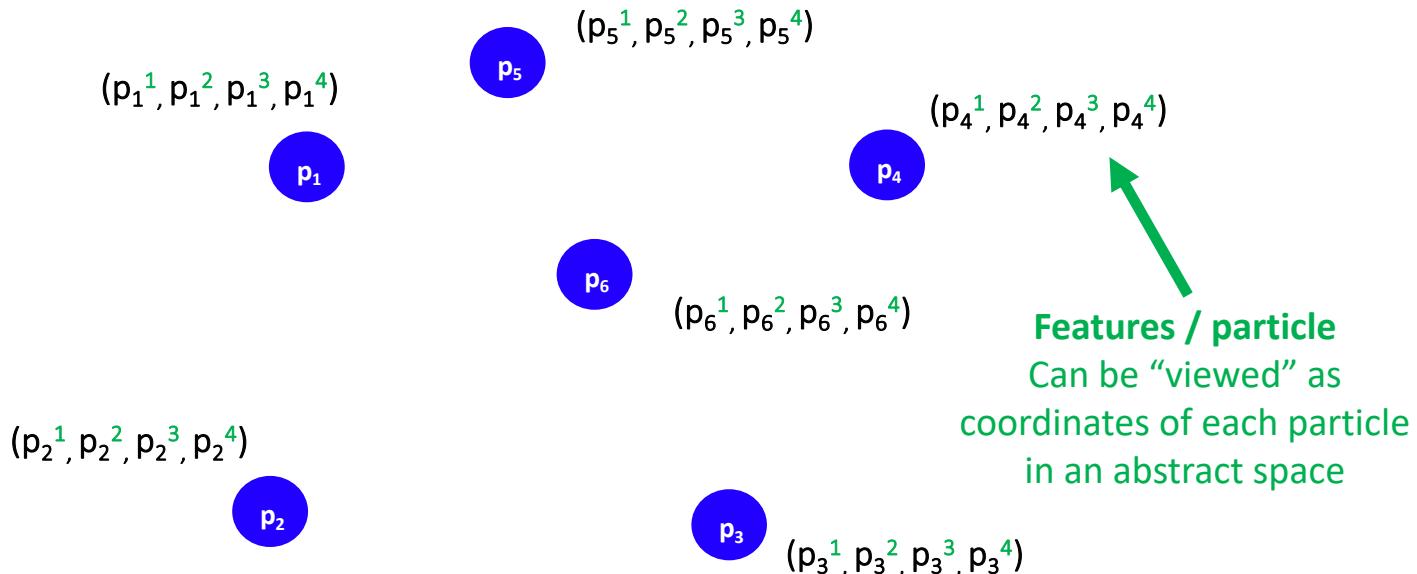
Dense NN

- Correlations & classification
- **Output:** binary or multiclass

Jet as a Graph

Jet as a Graph

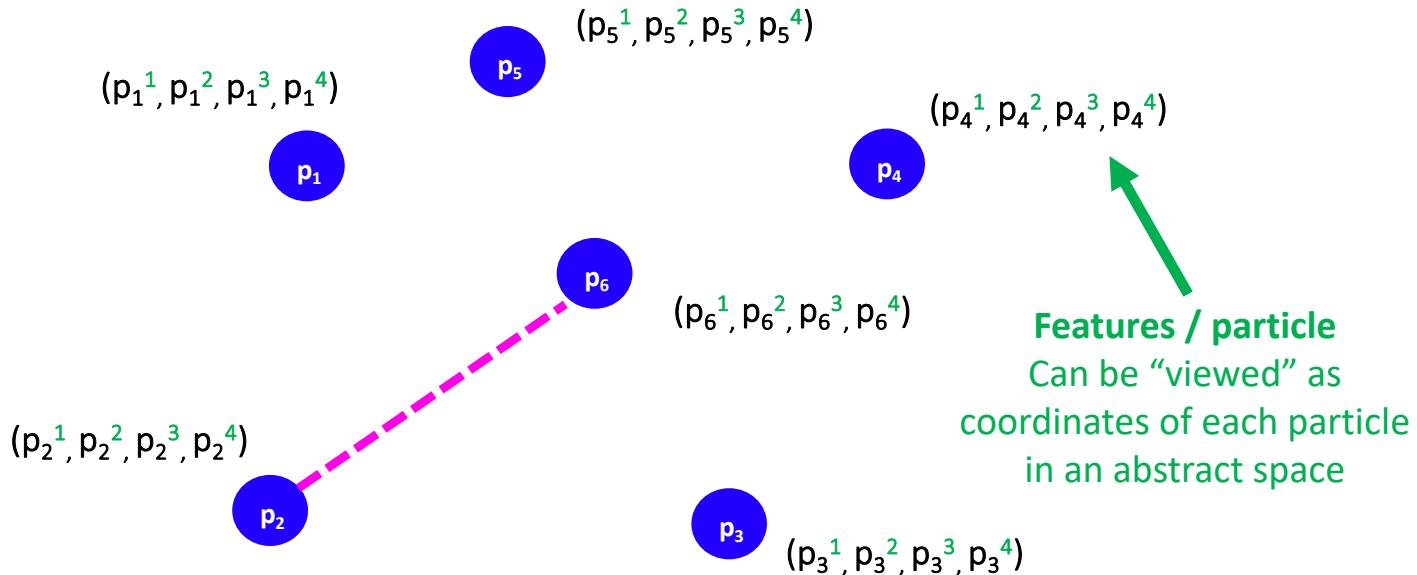
- We form a "Point/Particle Cloud"
 - ◆ Each particle: **node** of the graph



Jet as a Graph

- We form a "Point/Particle Cloud"

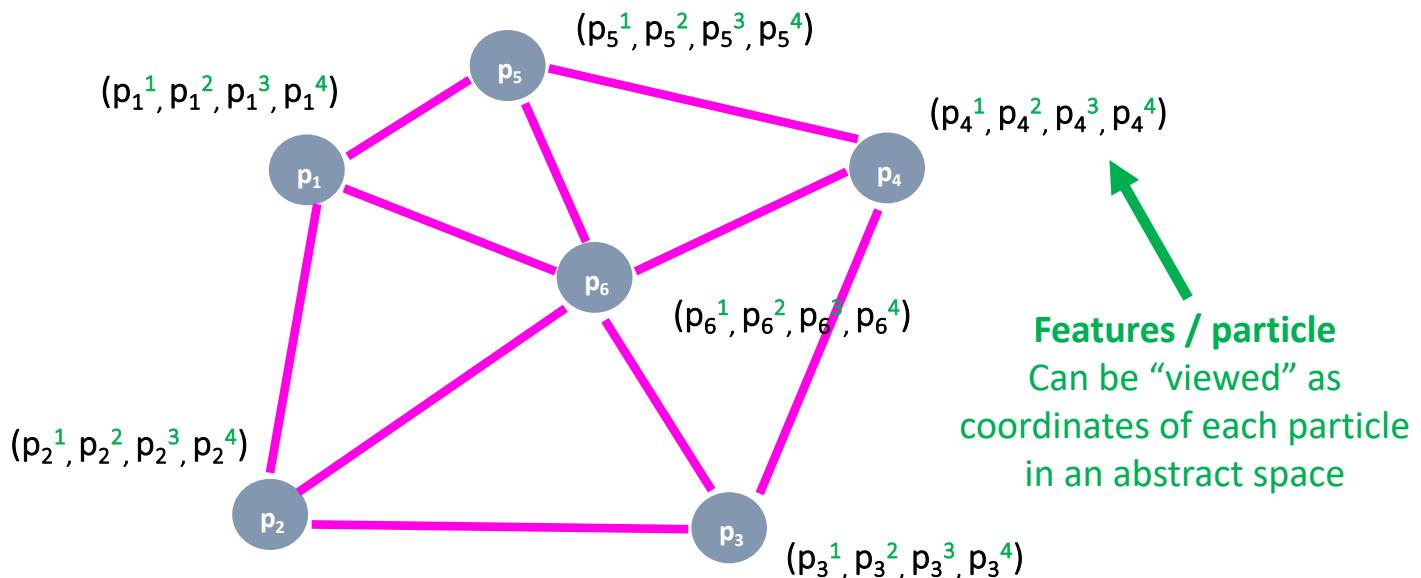
- ◆ Each particle: **node** of the graph



- Connections between particles: the **edges**

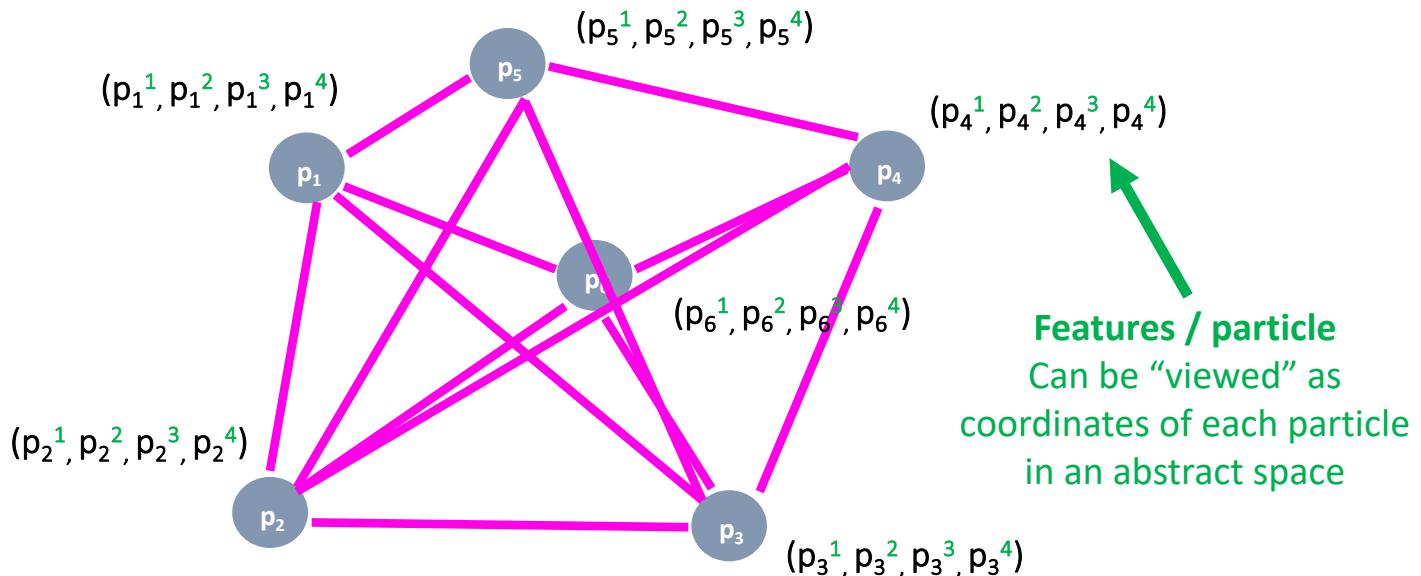
Step 1: Build the graph

- One approach: **Fully connected Graph**



Step 1: Build the graph

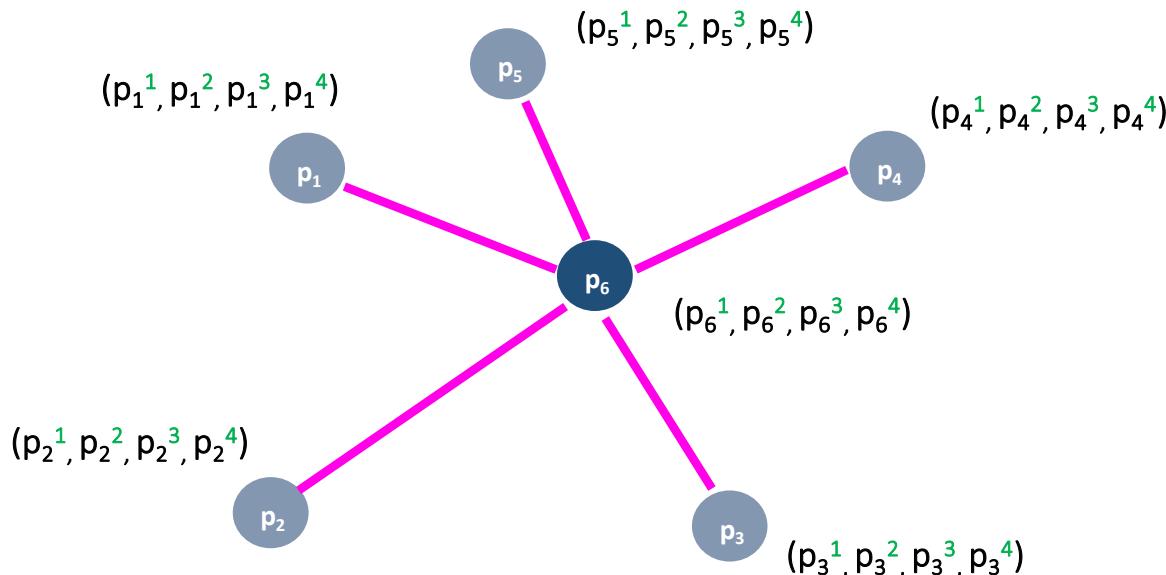
- One approach: **Fully connected Graph**



But computationally very expensive

Step 1: Build the graph

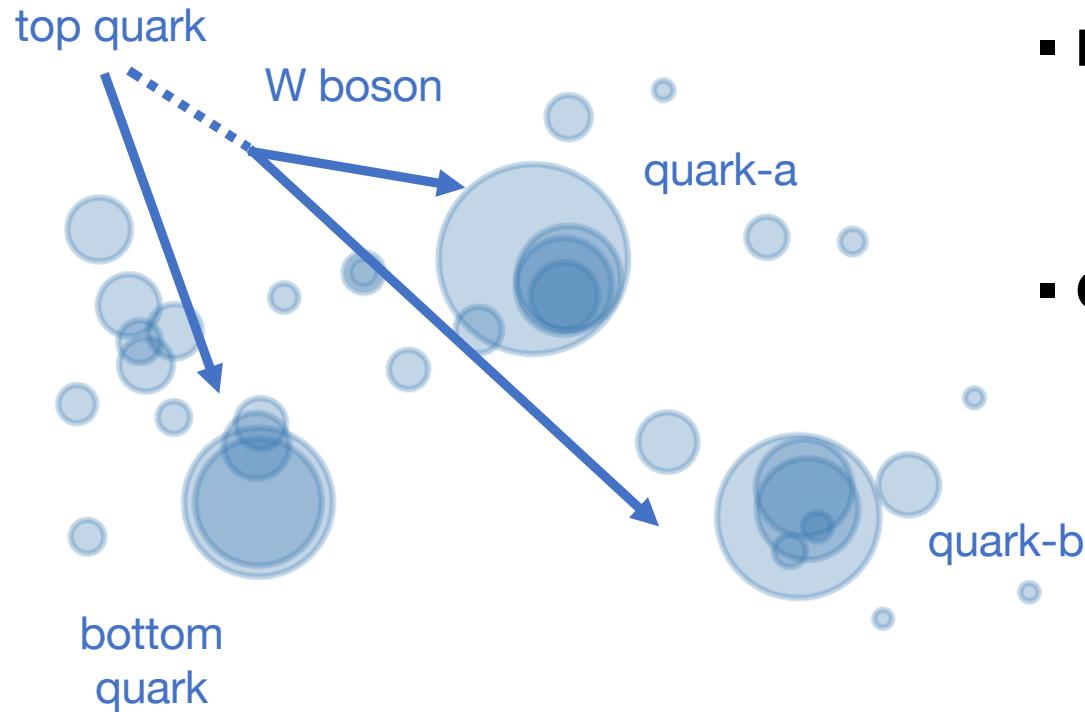
- One approach: **Fully connected Graph**
- **Another option:** apply some criteria
 - ◆ eg, select the k-Nearest Neighbors to each point, or within some distance, or



Step 2: Learn from the Graph

PRD 101 056019 (2020)

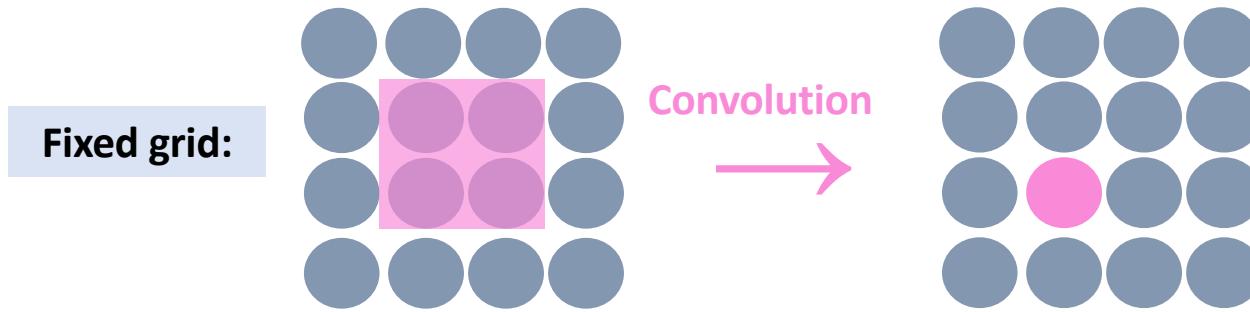
- Follow a **hierarchical learning** approach:
 - ◆ **First learn local structures and then more global ones**



- **Local** structures:
 - ◆ Identify the bottom quark and the two quarks from the W-boson separately
- **Global** info:
 - ◆ Presence of two quarks stemming from W-boson decay
 - ◆ Three quarks within jet, etc..

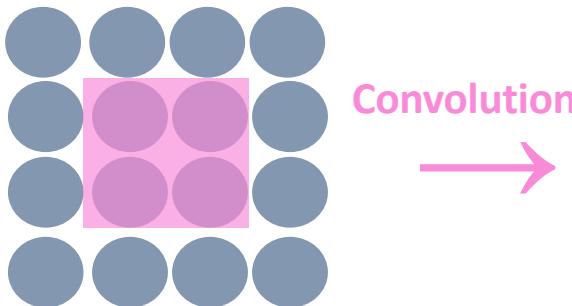
Convolution operations proven to be very powerful

Convolution on fixed grids [CNN-2D]

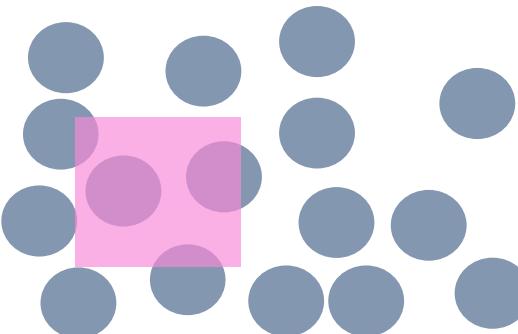


Convolution on Particle Clouds [?]

Fixed grid:



point/particle
cloud:



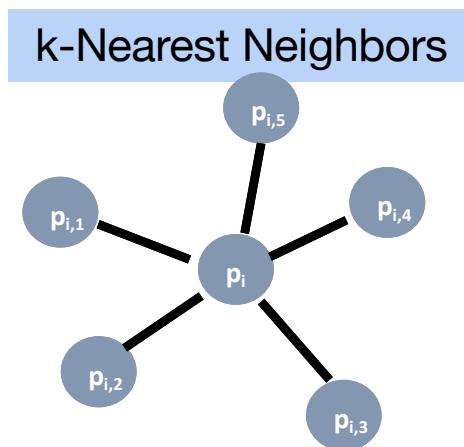
... but not straight forward on point/particle clouds

- Irregular and unordered sets
- Requires a permutation invariant convolutional operation

Convolution on particle clouds

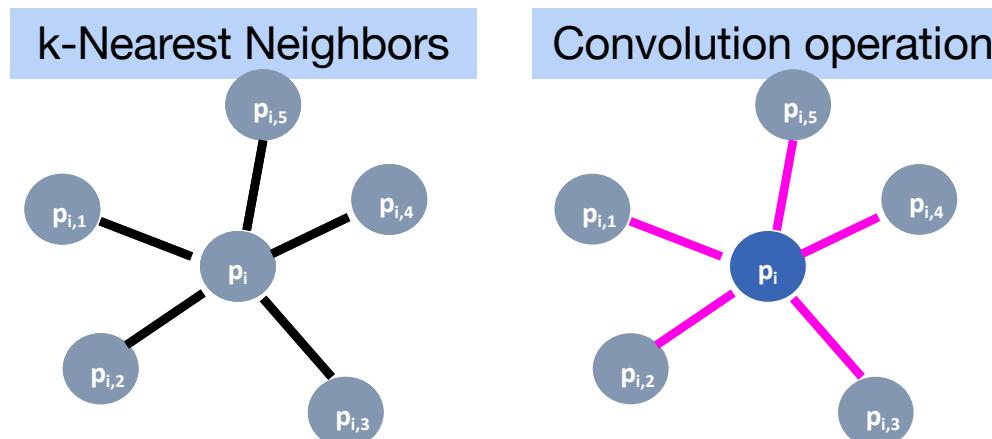
[arxiv:1801.07829](https://arxiv.org/abs/1801.07829)

- Find the **k -nearest neighbors** of each point



Convolution on particle clouds

- Find the **k -nearest neighbors** of each point
- Permutation invariant **convolution operation**
 - ◆ Edge feature func. → **aggregate** edge feat. w/ symmetric func.



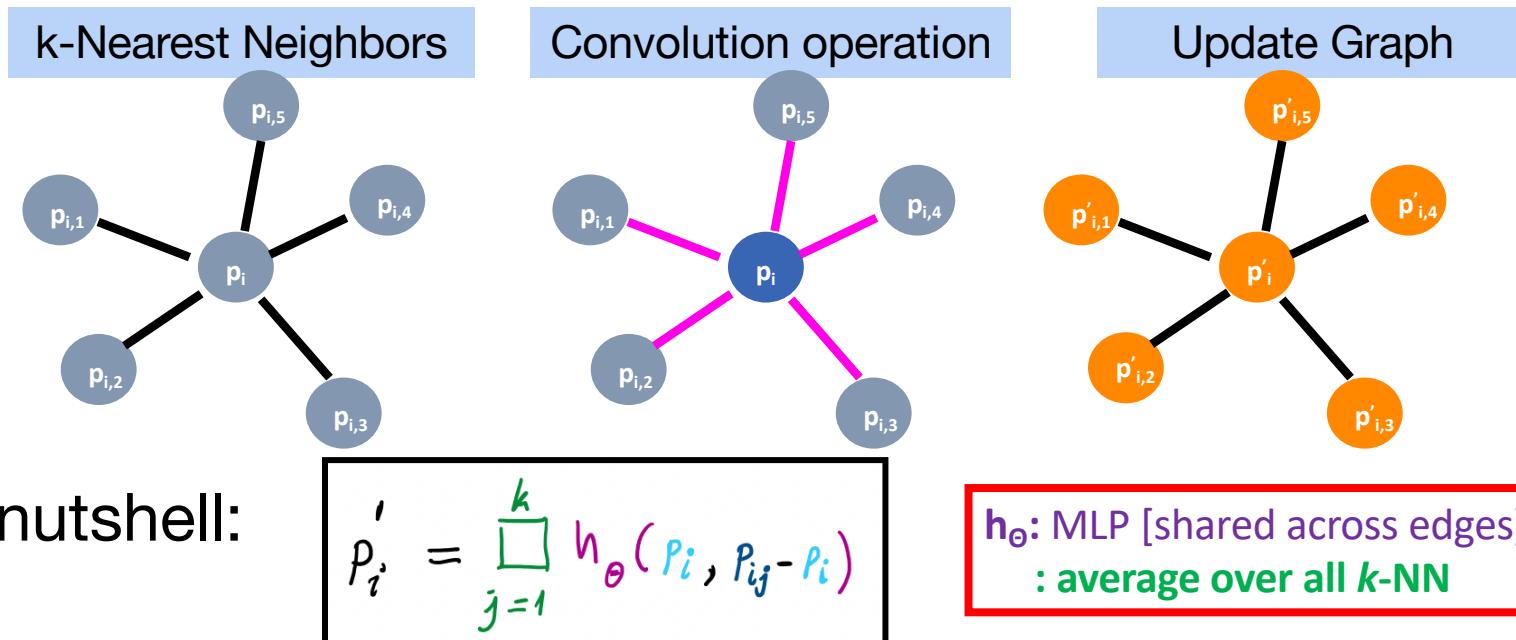
$$p'_i = \bigcup_{j=1}^k h_\theta(p_i, p_{ij} - p_i)$$

h_θ : MLP [shared across edges]
: average over all k -NN

- In a nutshell:

Convolution on particle clouds

- Find the **k -nearest neighbors** of each point
- Permutation invariant **convolution operation**
 - ◆ Edge feature func. → **aggregate** edge feat. w/ symmetric func.
- **Update Graph (ie Dynamic Graph CNN, DGCNN):**
 - ◆ Using k NN in the feature space produced after convolution

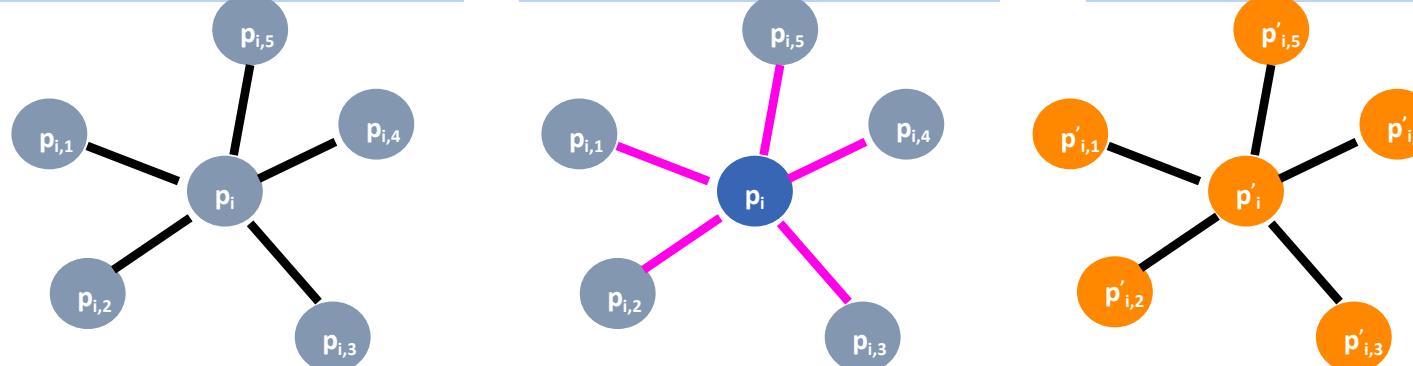


Convolution on particle clouds

A simplified version
during the hands-on
session ☺

- Find the **k -nearest neighbors** of each point
- Permutation invariant **convolution operation**
 - ◆ Edge feature func. → **aggregate** edge feat. w/ symmetric func.
- **Update Graph (ie Dynamic Graph CNN, DGCNN):**
 - ◆ Using k NN in the feature space produced after convolution

k-Nearest Neighbors Convolution operation Update Graph



- In a nutshell:

$$p'_i = \bigcup_{j=1}^k h_\theta(p_i, p_{ij} - p_i)$$

h_θ : MLP [shared across edges]
: average over all k -NN

Exercise 1: Hands-on part

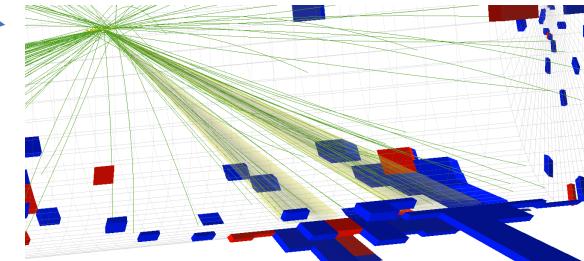
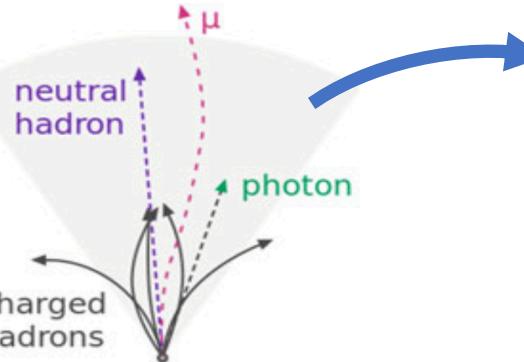
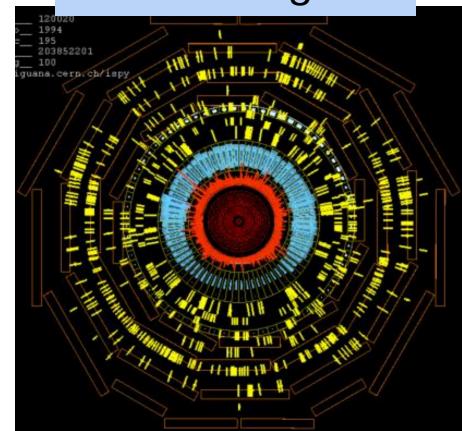
Exercise 2: Clustering

Particle/event reco chain

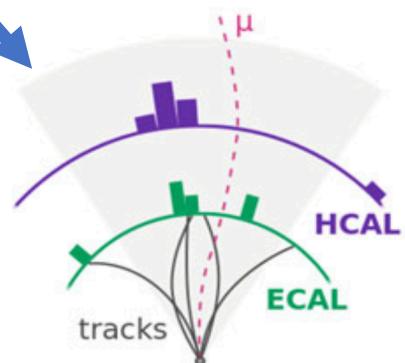
What we see:
Detector signals

Combine info from
different sub-detectors

Target:
Jet flavour



Output: γ , e, μ , charged & neutral hadrons



Cluster signals
in each sub-
detector
[Local RECO]

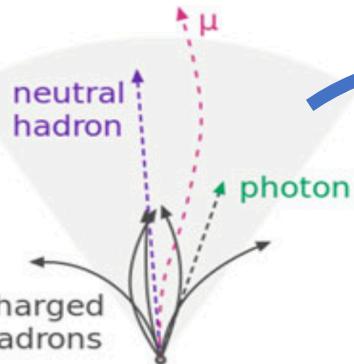
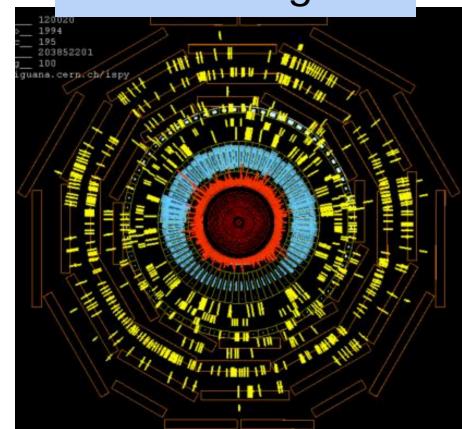
lower-level → **higher-level**

Particle/event reco chain

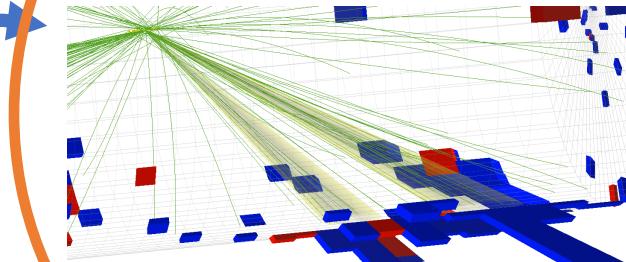
What we see:
Detector signals

Combine info from
different sub-detectors

Target:
Jet flavour

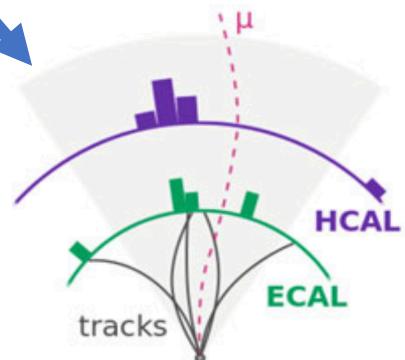


Output: γ , e, μ , charged &
neutral hadrons



Exercise 1

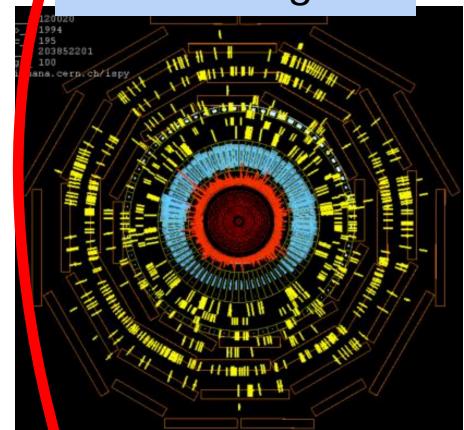
Cluster signals
in each sub-
detector
[Local RECO]



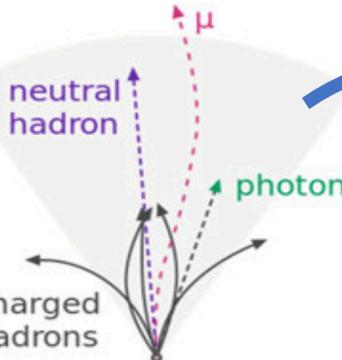
lower-level → **higher-level**

Particle/event reco chain

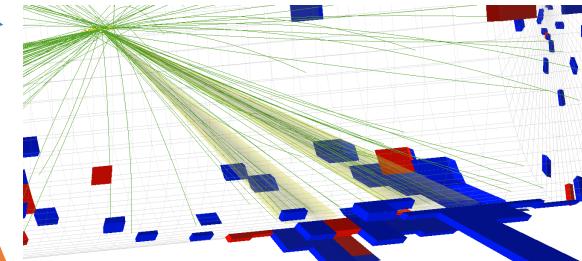
What we see:
Detector signals



Combine info from
different sub-detectors

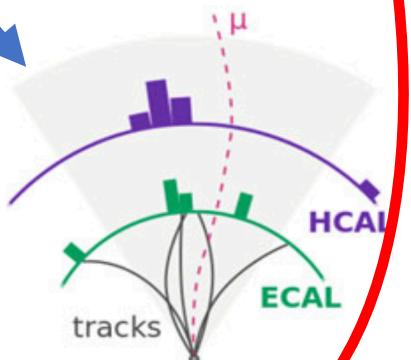


Target:
Jet flavour



Exercise 1

Output: γ , e, μ , charged &
neutral hadrons



Cluster signals
in each sub-
detector
[Local RECO]

Exercise 2

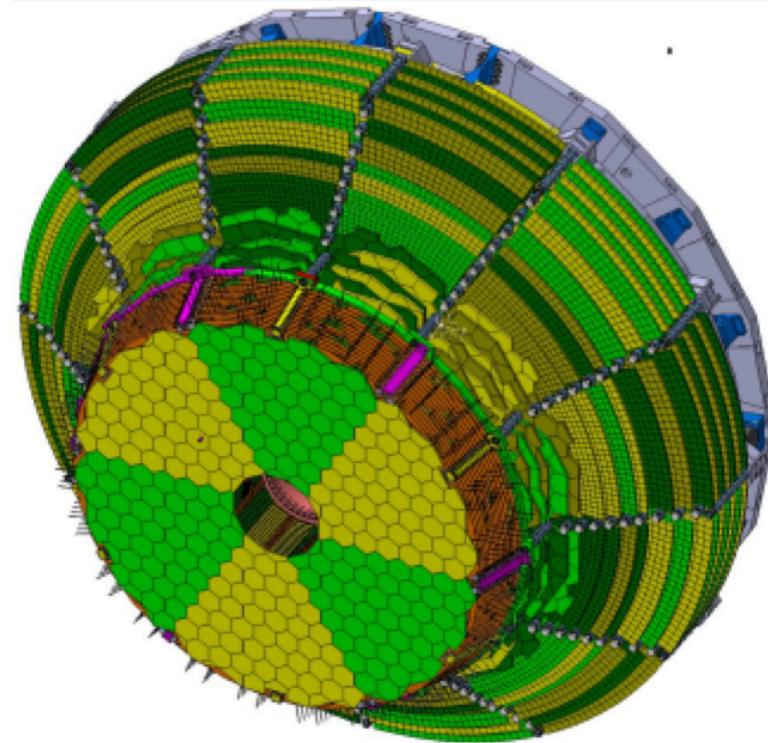
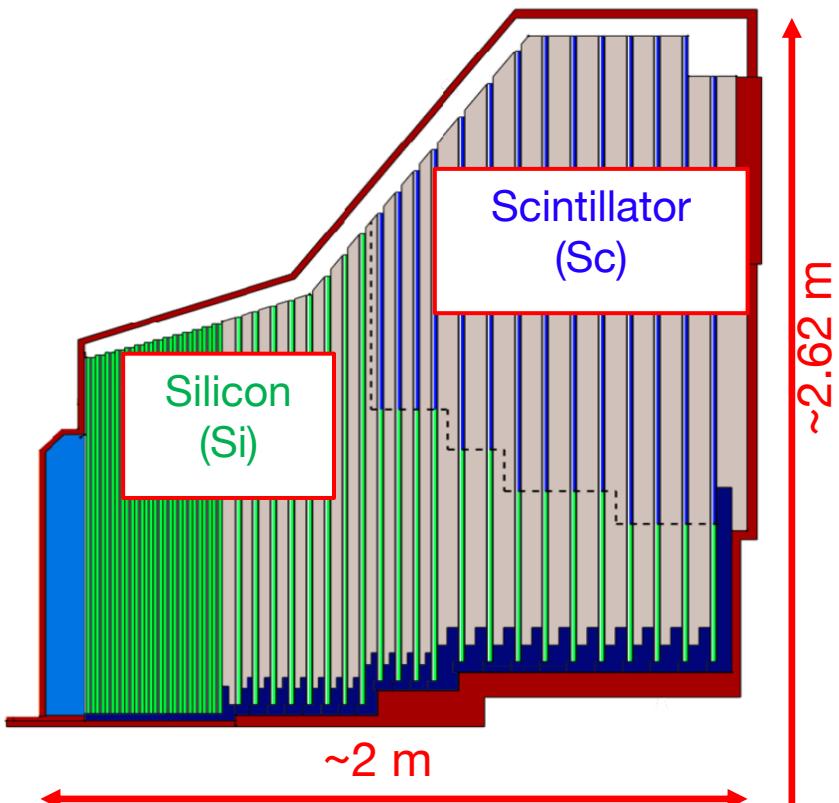
lower-level → higher-level

Exercise 2: Clustering

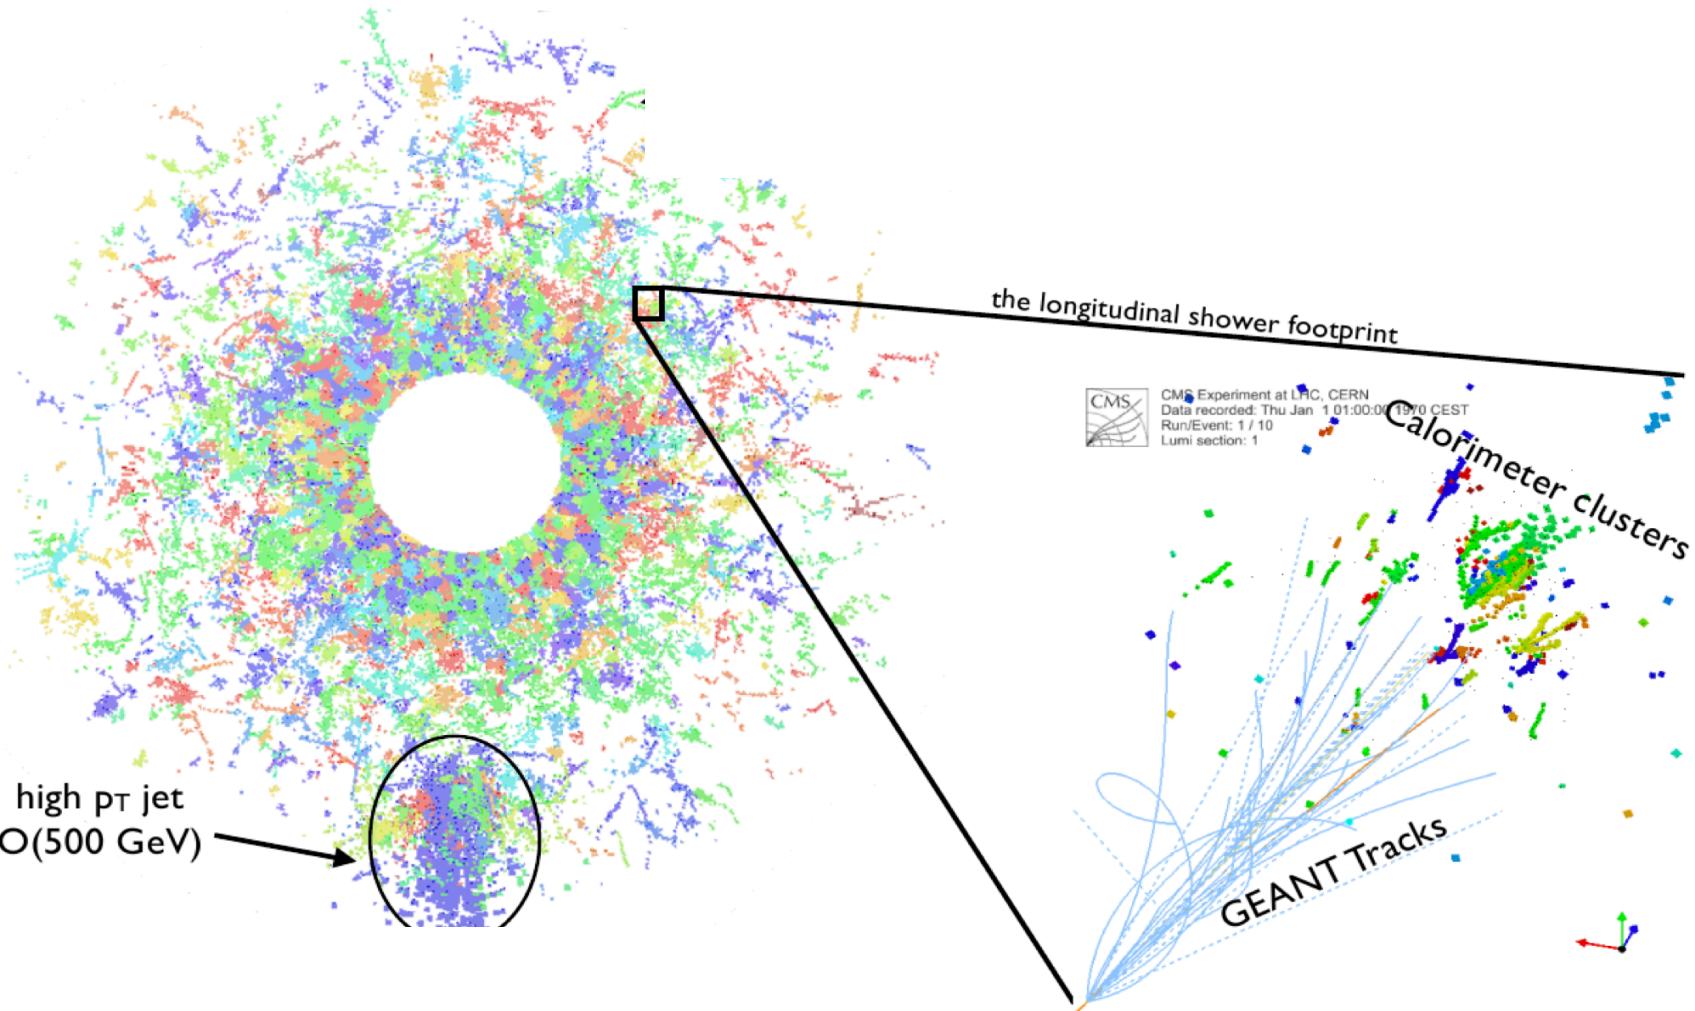
- Explore AI to cluster/group signals (e.g., energy deposits) in the detector
 - ◆ A particle traversing the detector give rise to signal in several channels
 - ◆ We aim to cluster deposits corresponding to each particle

Exercise 2: High Granularity Calorimeter

- Use a novel and quite ambitious detector project
 - ◆ The CMS High Granularity Endcap Calorimeter
 - Currently under construction; Target: 2030s



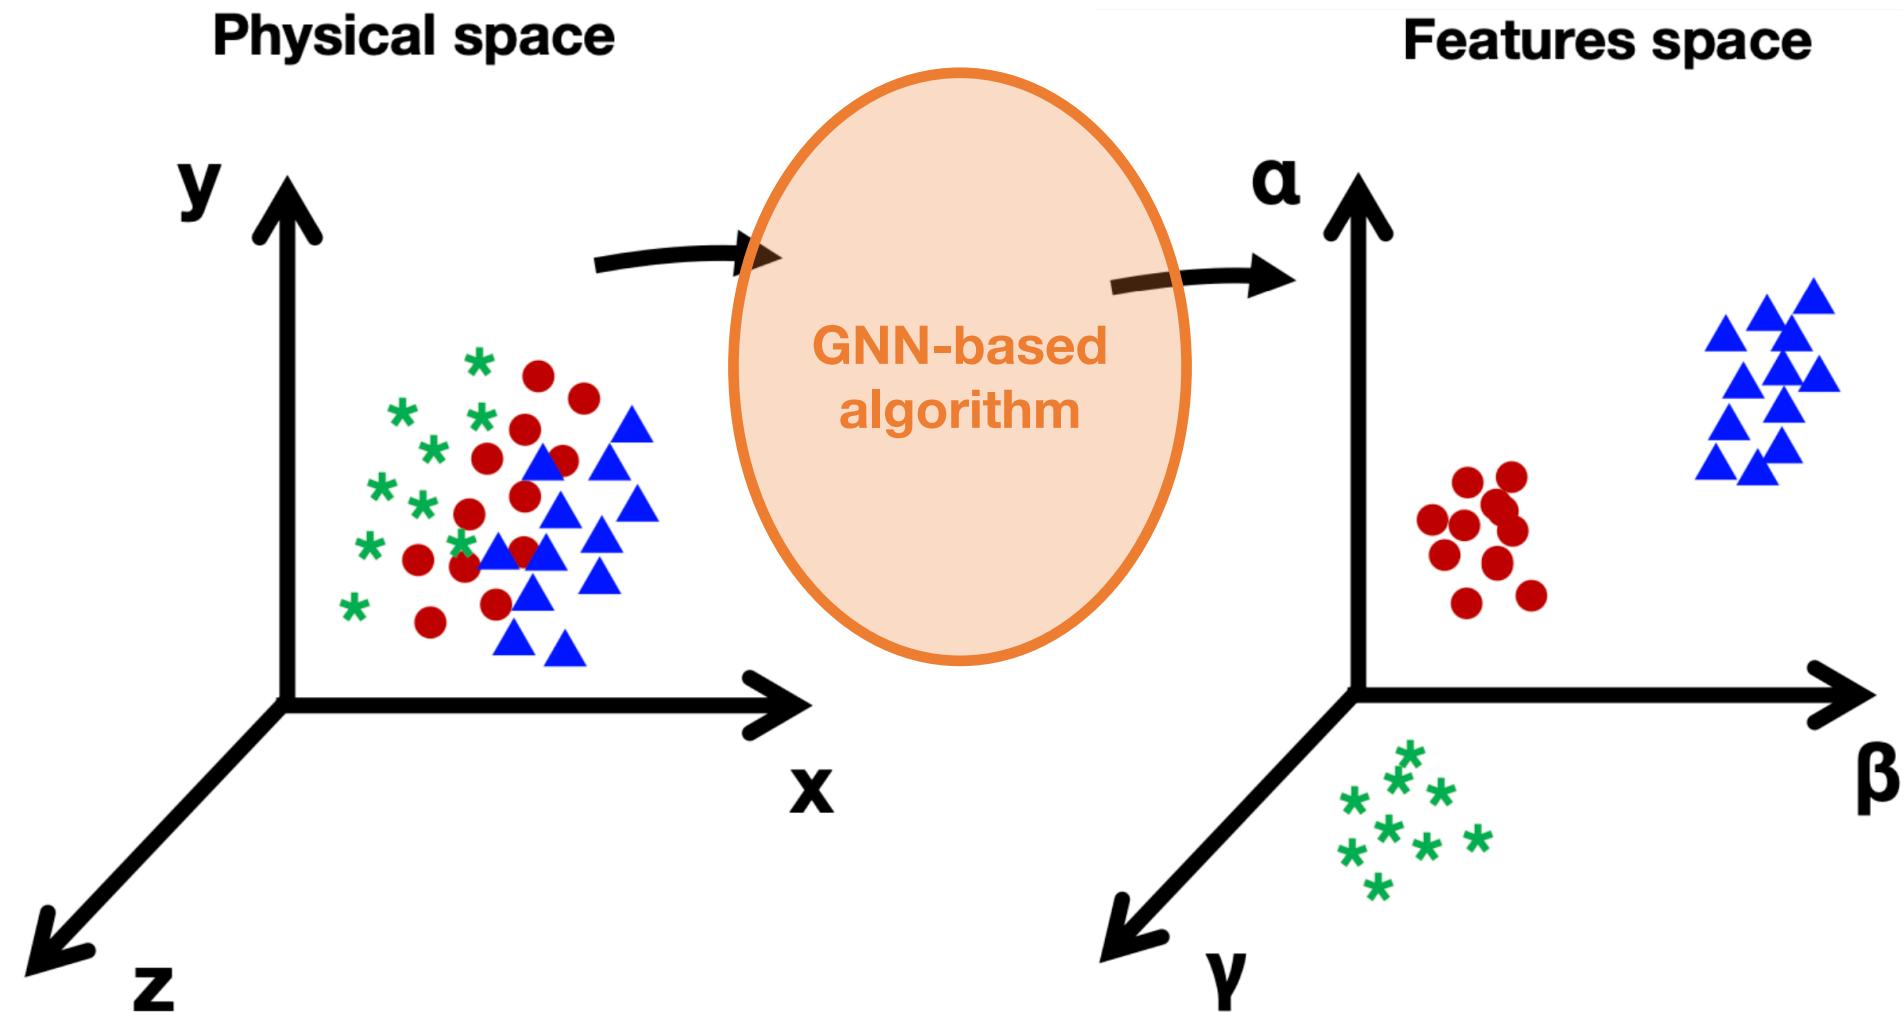
Simulated event



Clustering energy deposits

- Huge potential but extremely challenging using “traditional” (i.e., not ML) algorithms
 - ◆ each signal has a set of features: Energy, position, timing...
 - using these “raw” information is difficult to separate the deposits from different particles (i.e., different colors in the previous plot)
- GNN-based approach
 - ◆ Extract more discriminating set of features
 - in the latent/abstract space
 - ◆ Group them using a clustering algorithm
 - For simplicity: use k-Means in this exercise

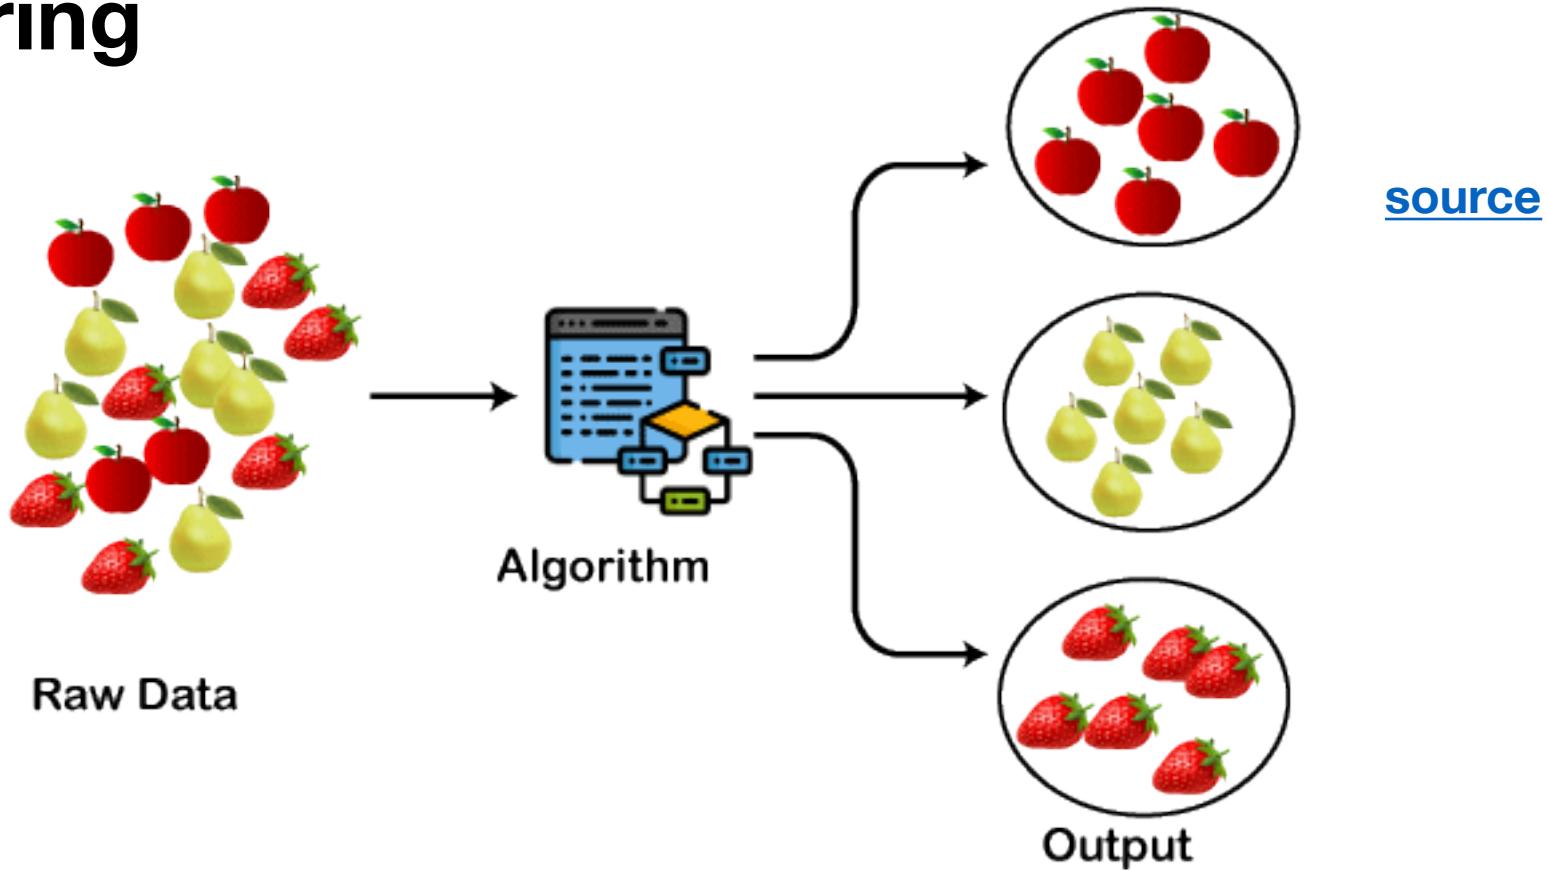
Clustering energy deposits



Clustering

- **Unsupervised** (or weakly-/self-supervised) learning method
 - ◆ Divides a set of points/entities/elements in groups with similar properties
- **Goal:** identify inherent structures/patterns within a dataset without using pre-labeled classes (i.e., no supervision).
 - ◆ Relies on a **similarity** measure
 - depends on the context and the data
- **Important:**
 - ◆ Since it is un-supervised learning we cannot use criteria such as accuracy, AUC, ..
 - challenging and highly subjective

Clustering



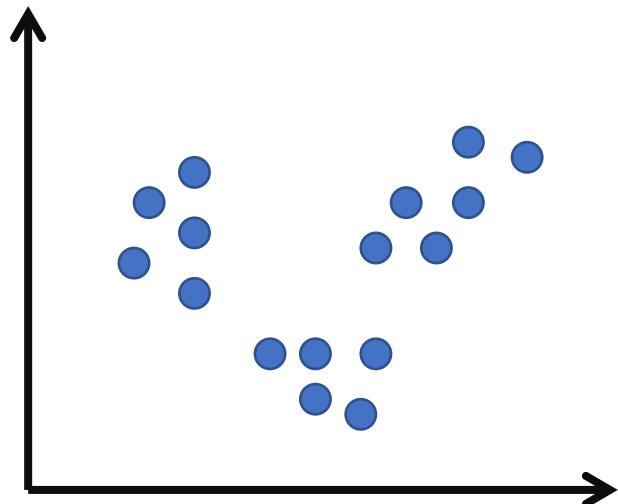
NB: No labelling [or very little]

Clustering algorithms aim to find patterns without us explicitly pointing them out.

Partitioning methods: k -Means

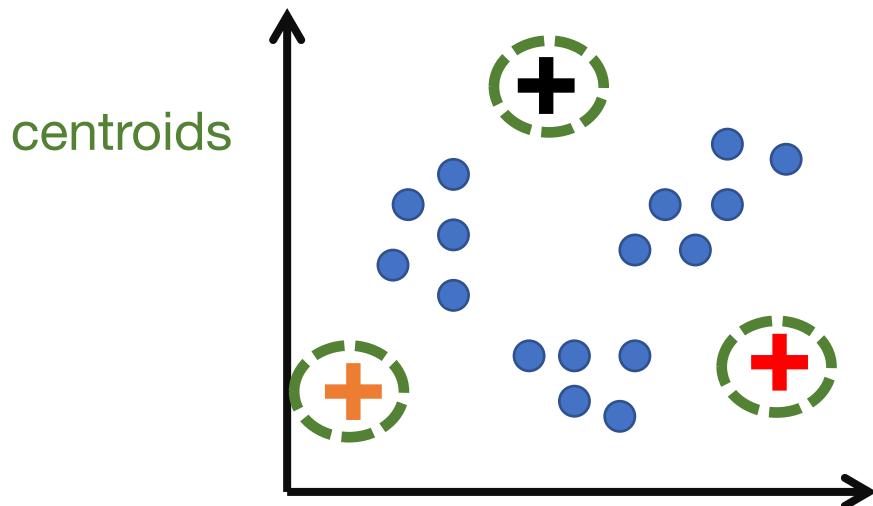
- Partition n observations into k clusters
 - ◆ each observation belongs to the cluster with the nearest mean
 - ◆ minimize an **objective function**

***k*-Means steps**



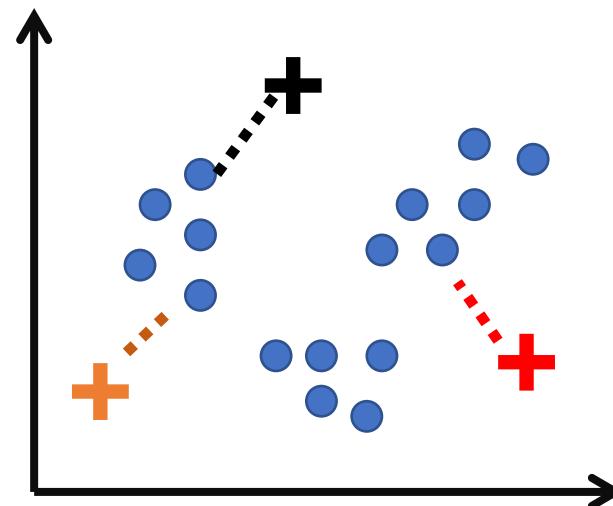
***k*-Means steps (I)**

- **Initialization:** Select ***k*** initial centroids [randomly or not]



k-Means steps (II)

- **Assignment:** Assign each data point to the nearest centroid, forming k-clusters:

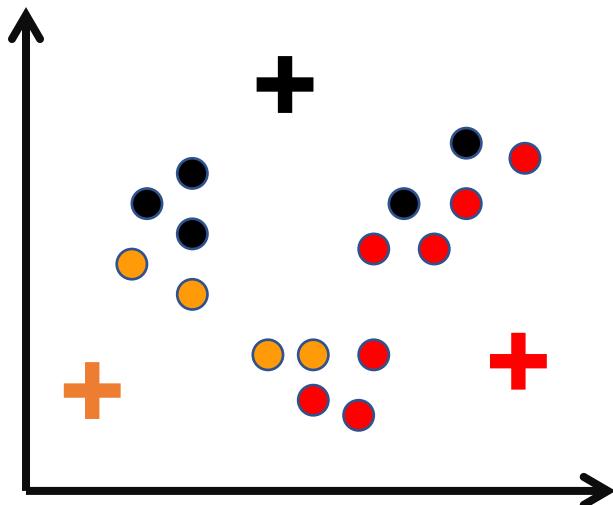


$$w_{ij} = \begin{cases} 1 & \text{if } j = \arg \min_{j'} \lVert x_i - \mu_{j'} \rVert^2 \\ 0 & \text{otherwise} \end{cases}$$

(Euclidian distance)²

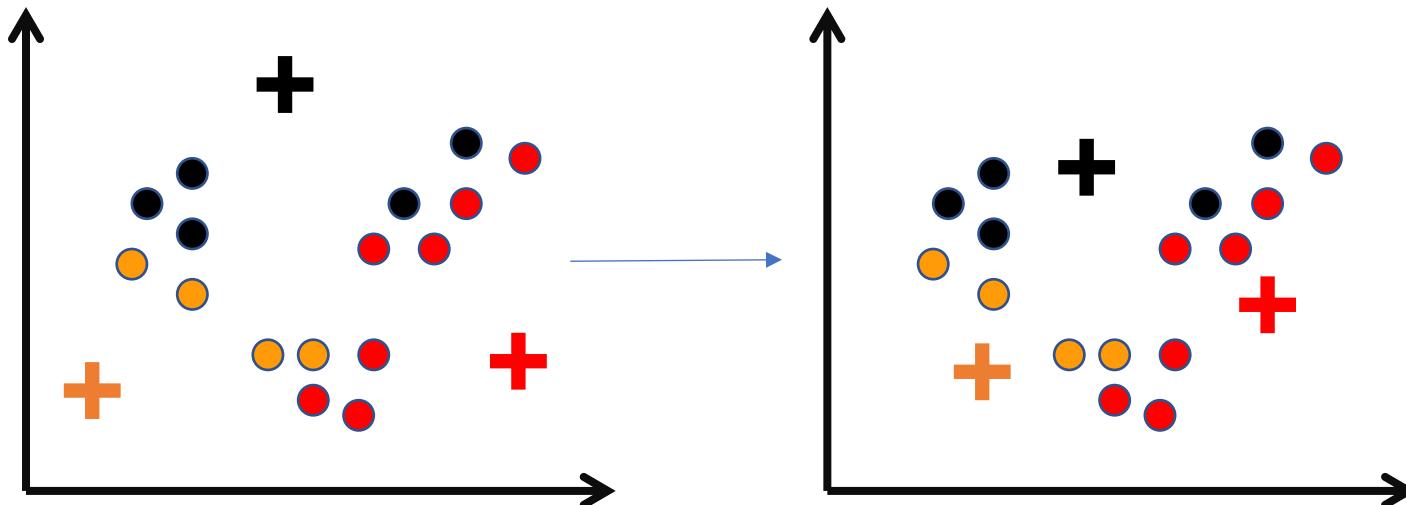
k-Means steps (III)

- **Update:** Recompute the centroid of each cluster as the mean of all points assigned to that cluster



k-Means steps (III)

- **Update:** Recompute the centroid of each cluster as the mean of all points assigned to that cluster



$$\mu_j = \frac{\sum_{i=1}^n w_{ij} x_i}{\sum_{i=1}^n w_{ij}}$$

k-Means steps (IV)

- **Repeat:** Assignment and Update steps repeated until centroids do not change significantly
 - ◆ minimize **objective function:**

$$J = \sum_{i=1}^n \sum_{j=1}^k w_{ij} \|x_i - \mu_j\|^2$$

1: if x_i assigned to cluster j
0: otherwise

k-Means steps (IV)

- **Repeat:** Assignment and Update steps repeated until centroids do not change significantly
 - ◆ minimize **objective function:**

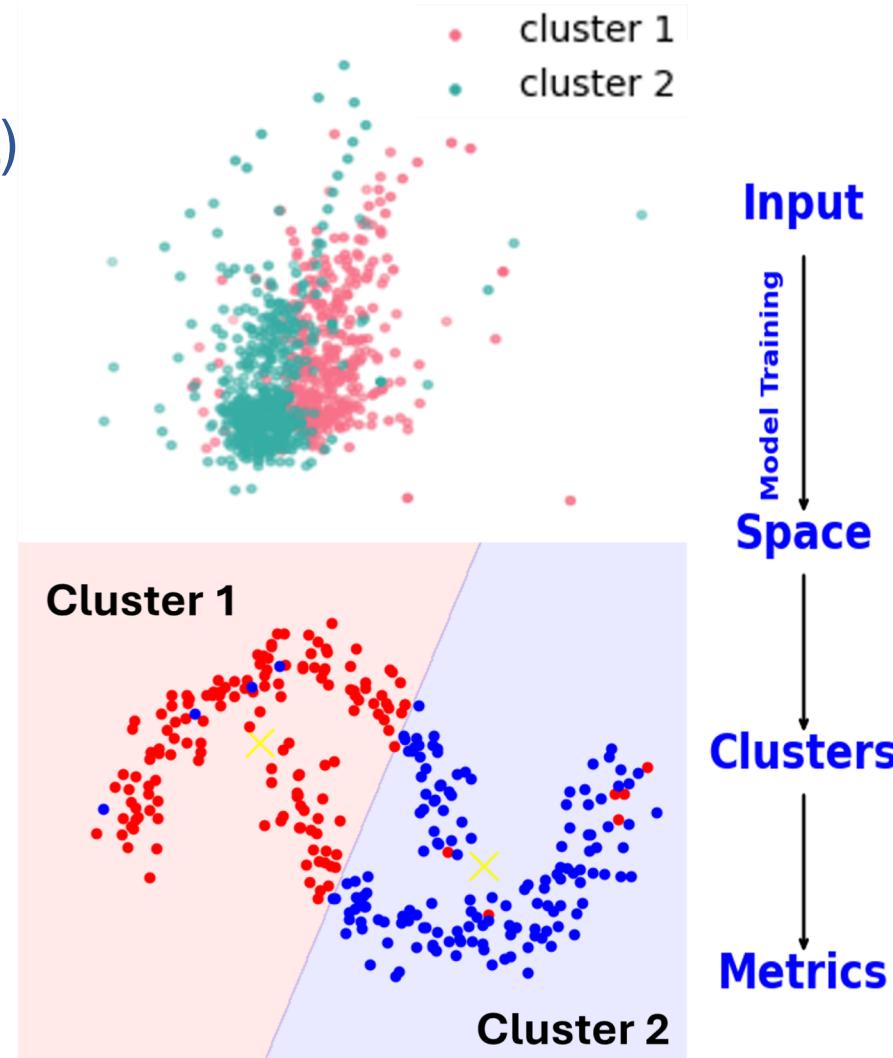
$$J = \sum_{i=1}^n \sum_{j=1}^k w_{ij} \|x_i - \mu_j\|^2$$

1: if x_i assigned to cluster j
0: otherwise

- Procedure stops when centroids no longer change significantly from one iteration to the next
 - ◆ algorithm has converged and objective function reached a local minimum

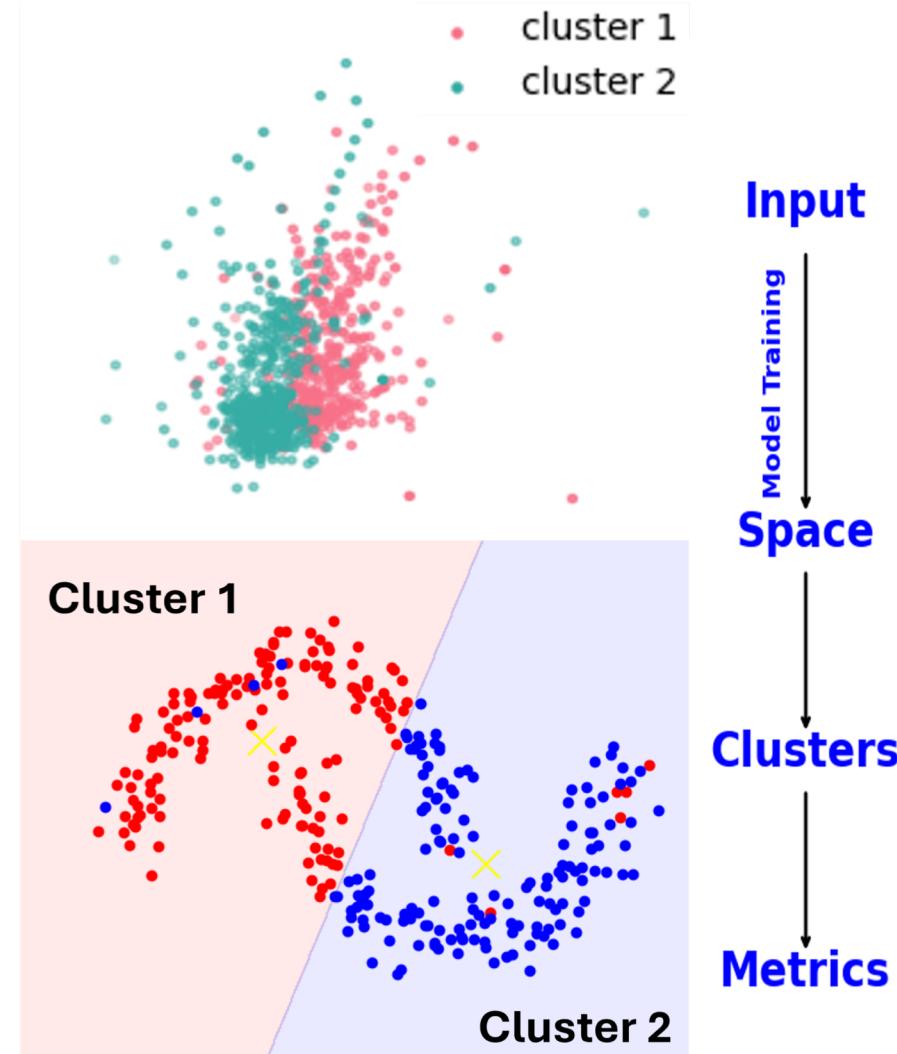
Hands-on 2: Algorithm Structure

- First:
 - Transform Input features (x,y,z,E) to new space through a GNN-based model
- Then:
 - Cluster with K-Means on model output



Hands-on 2: Alignment Diffusion Loss

- Minimize angle b/w connected vectors in the model's output
- Spread away angles with unconnected vectors



Hands-on 2: Metrics

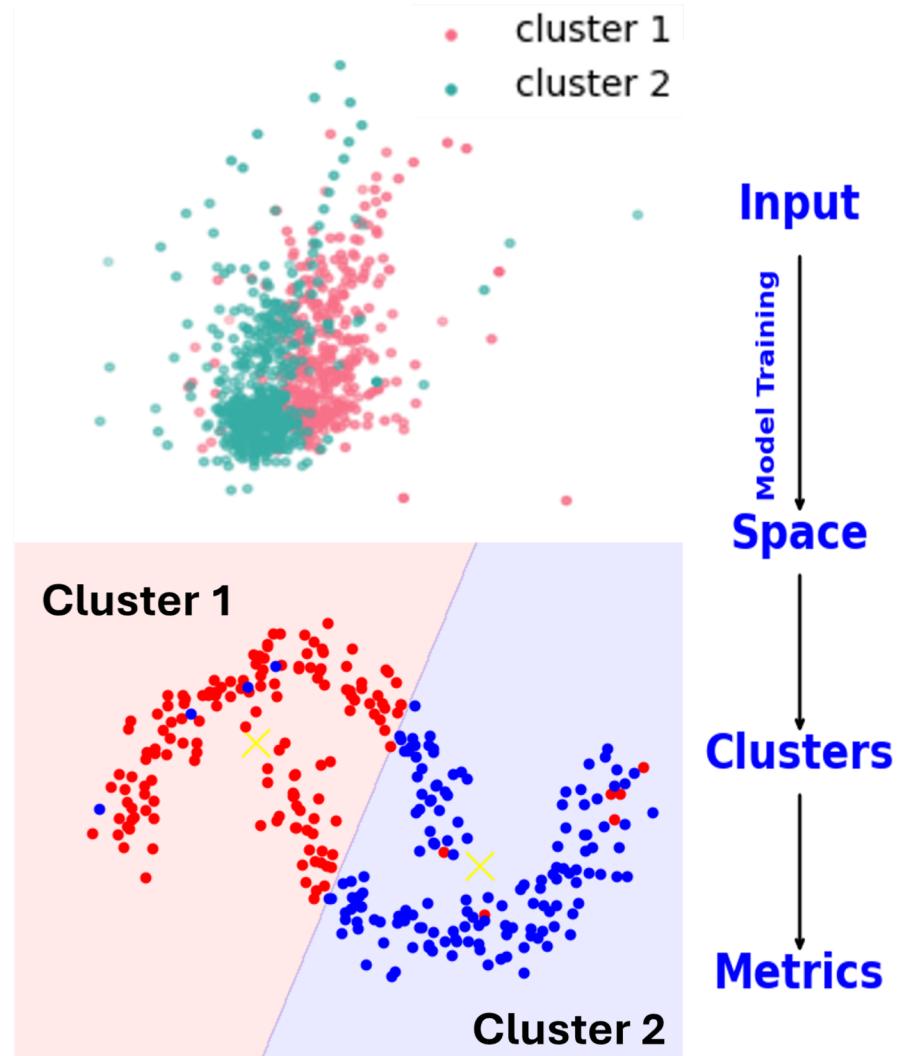
Evaluate how good the clusters are

Purity

Energy-weighted purity

$$p = \frac{\text{predicted}}{\text{actual}} \times 100\%$$

$$ewp = \frac{E_{\text{correct}}}{E_{\text{actual}}} \times 100\%$$



Exercise 2: hands-on